



جامعة الملك عبد الله  
للعلوم والتكنولوجيا  
King Abdullah University of  
Science and Technology

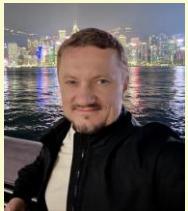
# Optimization Methods and Software for Federated Learning

PhD Defense

Konstantin Burlachenko  
Computer Science, KAUST

# Dissertation Defense Committee Members

#1



Peter Richtárik



Eric Feron



Suhaib Fahmy



Nic Lane



Stephen Boyd

Stanford  
University



UNIVERSITY OF  
CAMBRIDGE

Images: Google Search



# Background

## MS in Computer Science

Bauman Moscow State Technical University (2003 — 2009)

## Industry Experience

Startup (2012)      Acronis (2010 — 2012)      Yandex (2013 — 2014)

NVIDIA (2014 — 2019)      HUAWEI (2019 — 2020)



## Stanford Graduate Certificates

Data, Models and Optimization Graduate Certificate (2015 — 2018)

Artificial Intelligence Graduate Certificate (2016 — 2019)



## PhD Academic Journey

Joined Prof. P. Richtárik's Optimization and ML Lab at KAUST (August 2020)

Defended CS PhD Proposal (2022)

Member of Center of Excellence SDAIA-KAUST AI (2022 — 2023)

## Internships

Research Scientist Internship Offer, Facebook Inc., Menlo Park, USA (2021)

Internship in Private Federated Learning ML Team, Apple, Cambridge, UK (2024)

## Conference Presentations

Presentations: ICLR'24, SIAM'23, ICML'21, NSF-TRIPODS'21, DistributedML'21&'23

## Awards

Dean's Award (2020)

Grant from SDAIA (2022)

Dean's Award (2023)

AMD MI50 from AMD (2023)

Shaheen III Proposal (2024)

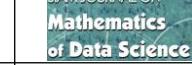
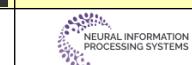
RDIA grant (2025)



Images: Google Search

# Papers Co-Authored During PhD (Creation Timeline)

#3

[1] Personalized Federated Learning with Communication Compression E. Bergou, <a href="#">K. Burlachenko</a> , A. Dutta, P. Richtárik	Ch6		 tmlr on ML RESEARCH
[2] MARINA: Faster Non-Convex Distributed Learning with Compression E. Gorbunov, <a href="#">K. Burlachenko</a> , Z. Li, P. Richtárik		Symposium on ACM PODC 2022	 ICML International Conference On Machine Learning
[3] FL_PyTorch: Optimization Research Simulator for Federated Learning <a href="#">K. Burlachenko</a> , S. Horváth, P. Richtárik	Ch2	Symposium on SIAM OP23	 acm  DISTRIBUTEDML
[4] Faster Rates for Compressed Federated Learning with Client-Variance Reduction H. Zhao, <a href="#">K. Burlachenko</a> , Z. Li, P. Richtárik			 SIAM JOURNAL ON Mathematics of Data Science
[5] Don't Compress Gradients in Random Reshuffling: Compress Gradient Differences A. Sadiev, G. Malinovsky, E. Gorbunov, I. Sokolov, A. Khaled, <a href="#">K. Burlachenko</a> , P. Richtárik		Workshop FL-ICML-2023	 NEURAL INFORMATION PROCESSING SYSTEMS
[6] Sharper Rates and Flexible Framework for Nonconvex SGD with Client and Data Sampling A. Tyurin, L. Sun, <a href="#">K. Burlachenko</a> , P. Richtárik	Ch5		 tmlr on ML RESEARCH
[7] Federated Learning with Regularized Client Participation G. Malinovsky, S. Horváth, <a href="#">K. Burlachenko</a> , P. Richtárik		Workshop FL-ICML-2023	 arXiv UNDER REVIEW
[8] Error Feedback Shines when Features are Rare P. Richtárik, E. Gasanov, <a href="#">K. Burlachenko</a>			 arXiv Preparing for Resubmission
[9] Federated Learning is Better with Non-Homomorphic Encryption <a href="#">K. Burlachenko</a> , A. Alrowithi, F. Ali Albalawi, P. Richtárik	Ch4		 acm  DISTRIBUTEDML
[10] Error Feedback Reloaded: From Quadratic to Arithmetic Mean of Smoothness Constants P. Richtárik, E. Gasanov, <a href="#">K. Burlachenko</a>	Ch3	ML Summer School Okinawa 2024	 ICLR International Conference On Learning Representations
[11] Unlocking FedNL: Self-Contained Compute-Optimized Implementation <a href="#">K. Burlachenko</a> , P. Richtárik	Ch7	KAUST AI Symposium 2024	 arXiv UNDER REVIEW
[12] PV-Tuning: Beyond Straight-Through Estimation for Extreme LLM Compression V. Malinovskii, D. Mazur, I. Ilin, D. Kuznedelev, <a href="#">K. Burlachenko</a> , K. Yi, D. Alistarh, P. Richtárik			 NEURAL INFORMATION PROCESSING SYSTEMS
[13] BurTorch: Revisiting Training from First Principles by Coupling Autodiff, Math Optimization, and Systems <a href="#">K. Burlachenko</a> , P. Richtárik	Ch8		 arXiv UNDER REVIEW

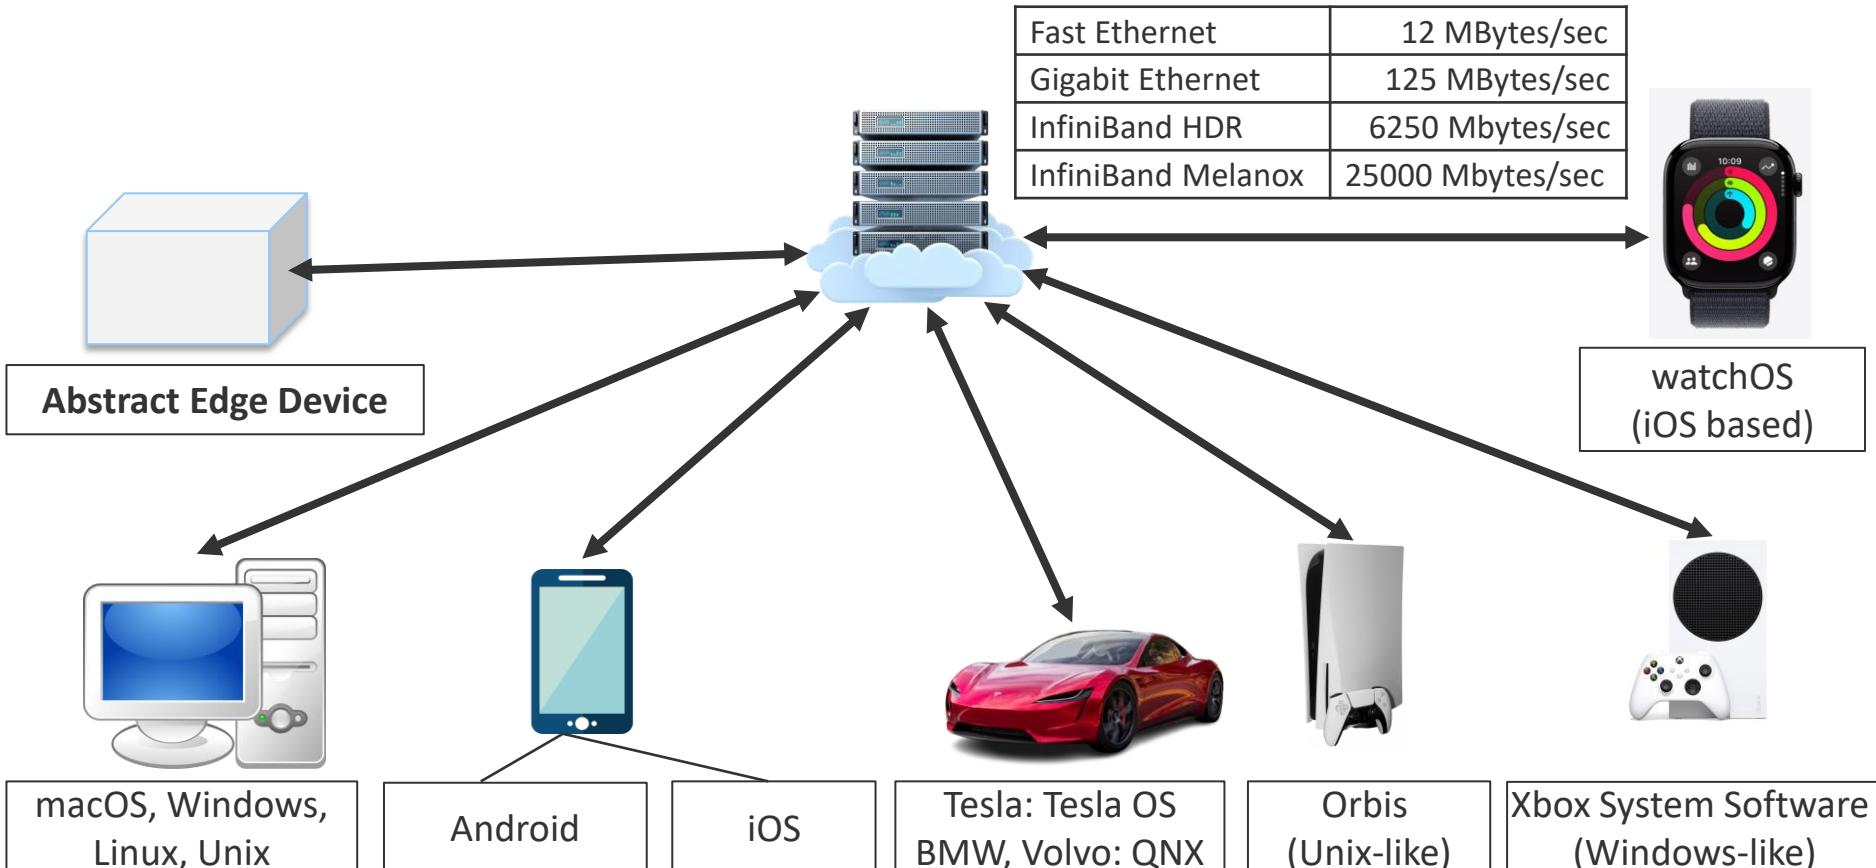
Traditional machine learning assumes that  
the training dataset is collected and stored centrally

Traditional machine learning assumes that  
the training dataset is collected and stored centrally

However, centralized storage is **not** where data is generated in the first place

# Shifting Training to Edge Devices

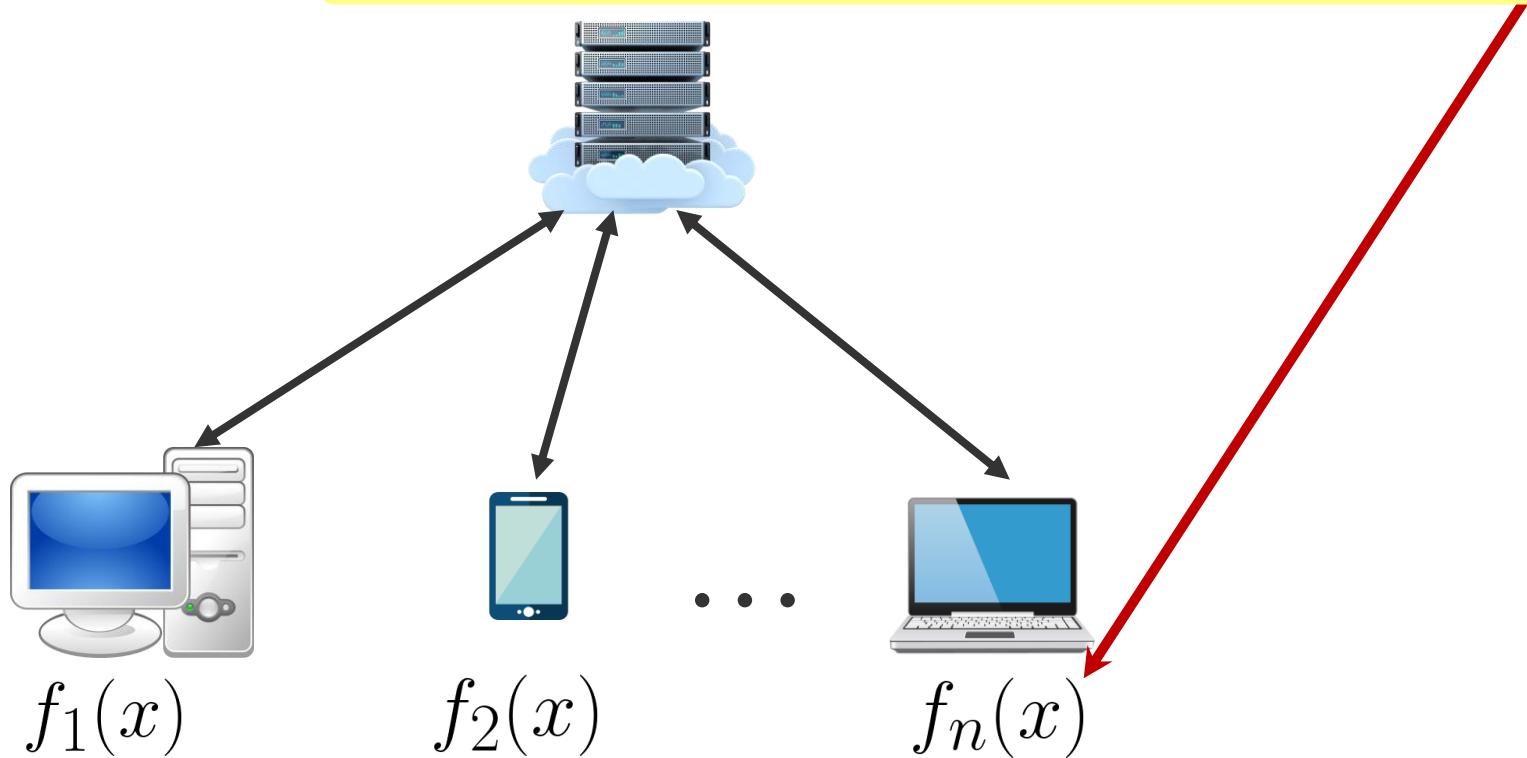
#5



Images: Google Search

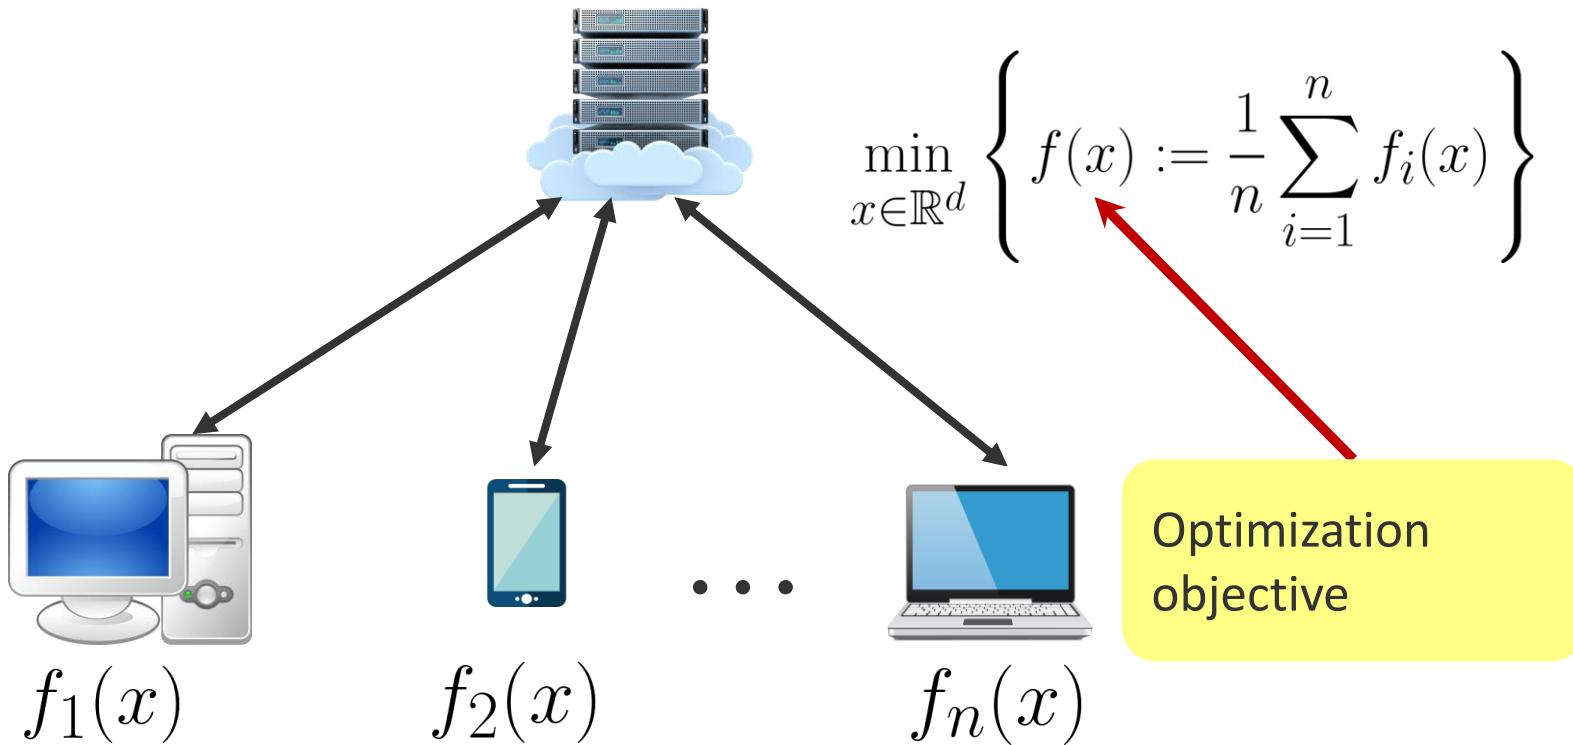
# Shifting Training to Edge Devices

$f_n$  is a local loss constructed from data  $D_n$



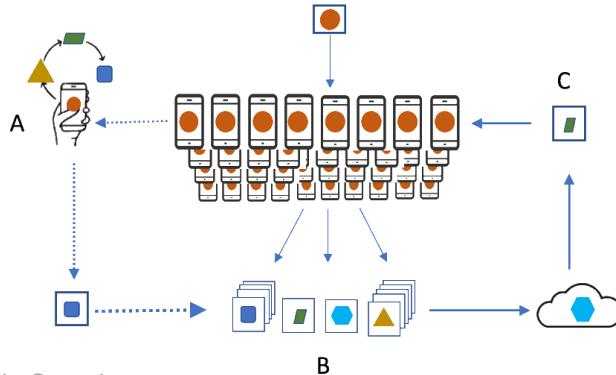
# Shifting Training to Edge Devices

#6





Images: Google Search



$$\min_{x \in \mathbb{R}^d} \left\{ f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x) \right\}$$

## FL Origins

Federated Learning: Strategies for Improving Communication Efficiency (2016) J. Konečný, B. McMahan, F. X. Yu, P. Richtárik, A.T. Suresh, D. Bacon  
Federated Optimization: Distributed Machine Learning for On-Device Intelligence (2016) J. Konečný, B. McMahan, D. Ramage, P. Richtárik  
Communication-Efficient Learning of Deep Networks from Decentralized Data (2017) B. McMahan, et al.  
Advances and Open Problems in Federated Learning (2021) P. Kairouz, et al.

The first publication with “Federated Learning” in its title

While FL mitigates sample size limitations and enables novel decentralized applications, it also brings new challenges

# Federated Learning Challenges Addressed in the Thesis #8

## Theoretical Work

1. Data Heterogeneity

2. Device Heterogeneity

3. Communication Bottleneck

4. Privacy

5. Software

## Theory-Inspired Practical Work

## Practical Work

### Ch1: Introduction

Ch3: EF21-W  
Richtárik et al., 2024

Ch3: EF21-W  
Richtárik et al., 2024

Ch4: DCGD/PERMK/AES  
Burlachenko et al., 2023

Ch5: PAGE Extensions  
Tyurin et al., 2023

Ch6:  
Compressed L2GD  
Bergou et al., 2023

Ch7: Unlocking FedNL  
Burlachenko and Richtárik, 2024

Ch8: BurTorch  
Burlachenko & Richtárik, 2025

Ch2: FL\_PyTorch  
Burlachenko et al., 2021

Ch9: Concluding Remarks: Summary and Future Research

# Federated Learning Challenges Addressed in the Thesis #9

## Theoretical Work

1. Data Heterogeneity

2. Device Heterogeneity

3. Communication  
Budget Check

4. Privacy

5. Software

## Theory-Inspired Practical Work

## Practical Work

Ch1: Introduction

Ch3: EF21-W  
Richtárik et al., 2024

Ch3: EF21-W  
Richtárik et al., 2024

Ch4: DCGD/PERMK/AES  
Burlachenko et al., 2023

Ch5: PAGE Extensions  
Tyurin et al., 2023

Ch6:  
Compressed L2GD  
Bergou et al., 2023

Ch7: Unlocking FedNL  
Burlachenko and Richtárik, 2024

Ch8: BurTorch  
Burlachenko & Richtárik, 2025

Ch2: FL\_PyTorch  
Burlachenko et al., 2021

Ch9: Concluding Remarks: Summary and Future Research

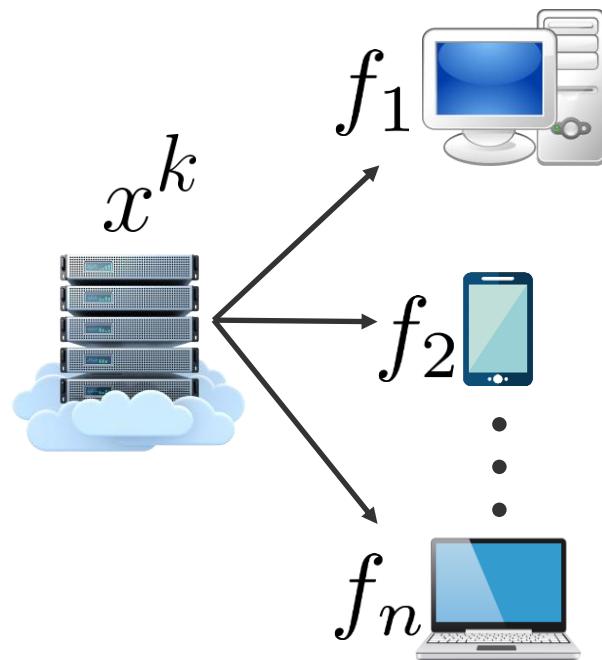
# Distributed Gradient Descent

#10


$$x^k$$


# Distributed Gradient Descent

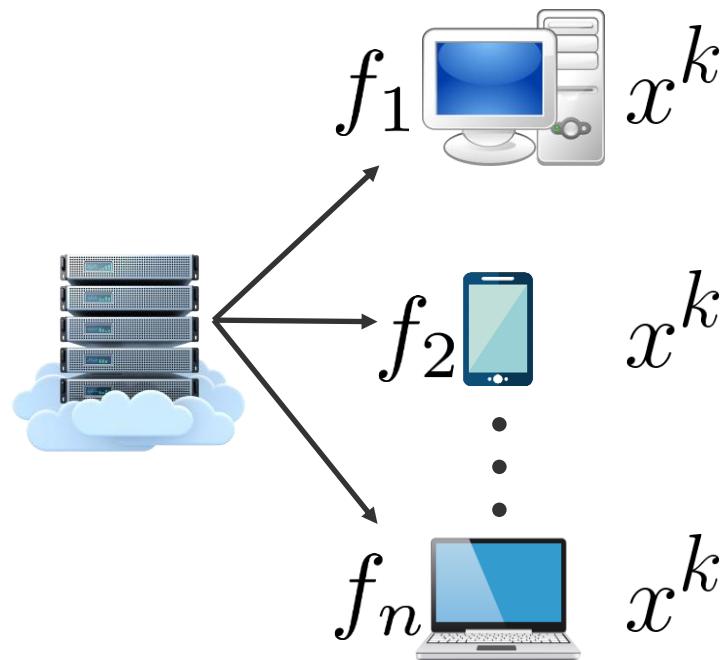
#10



# Distributed Gradient Descent

#10

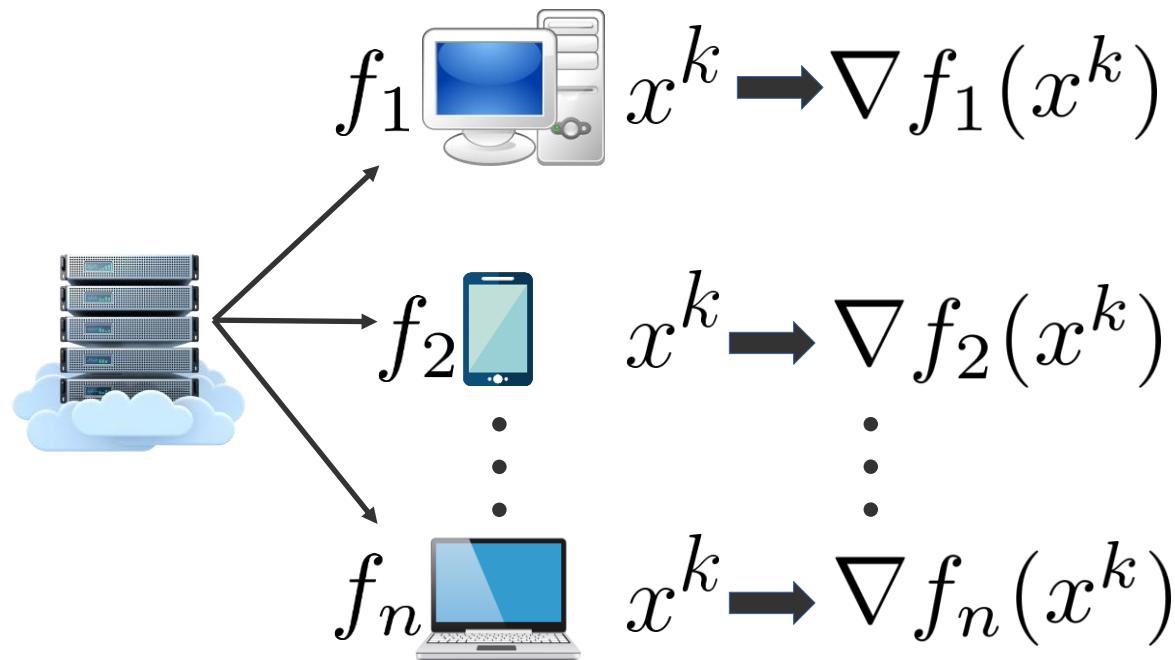
○ ○ ● ○ ○ ○



# Distributed Gradient Descent

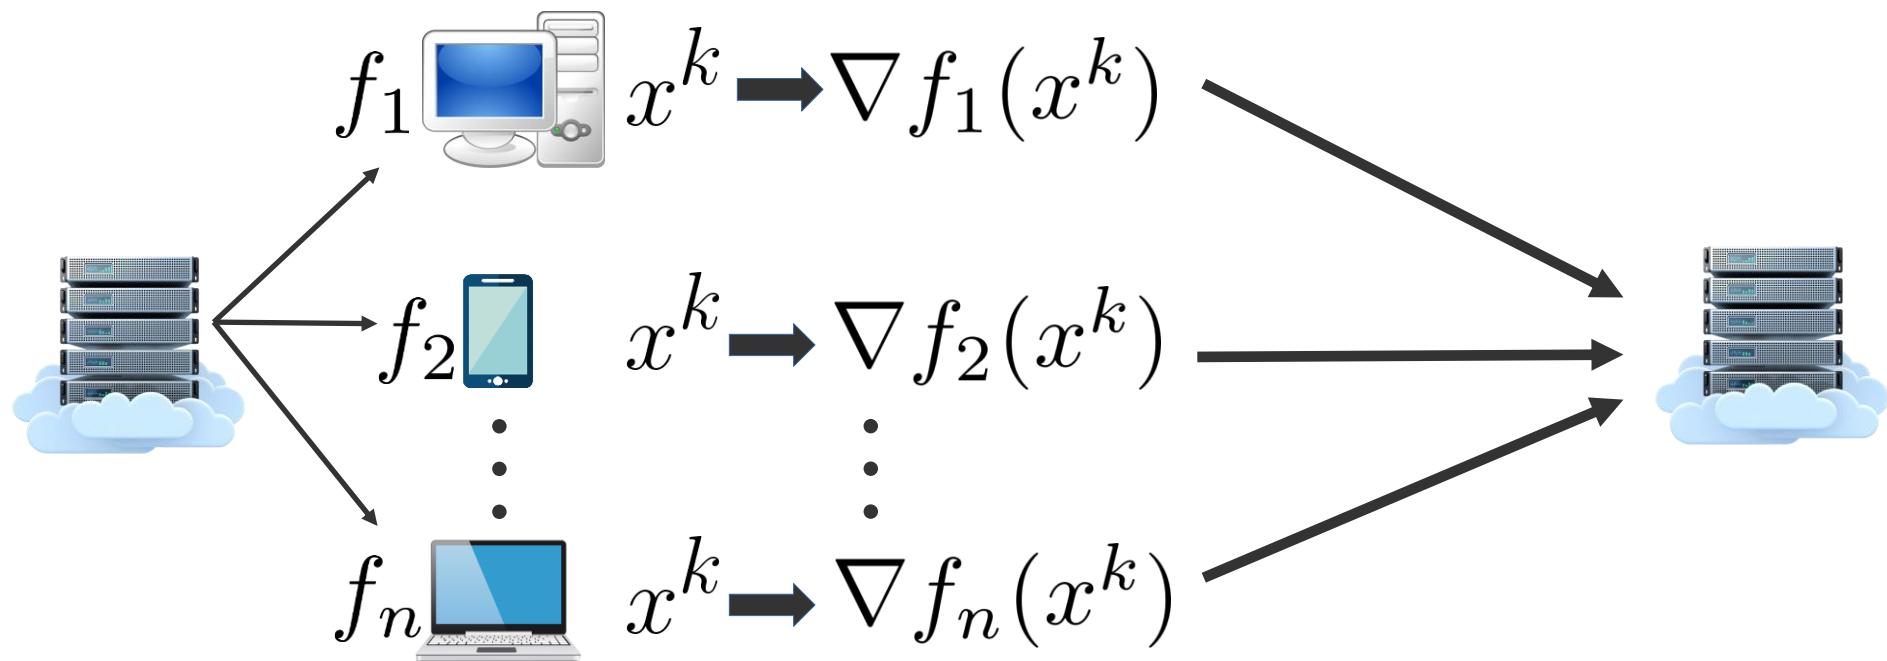
#10

○ ○ ○ ● ○ ○



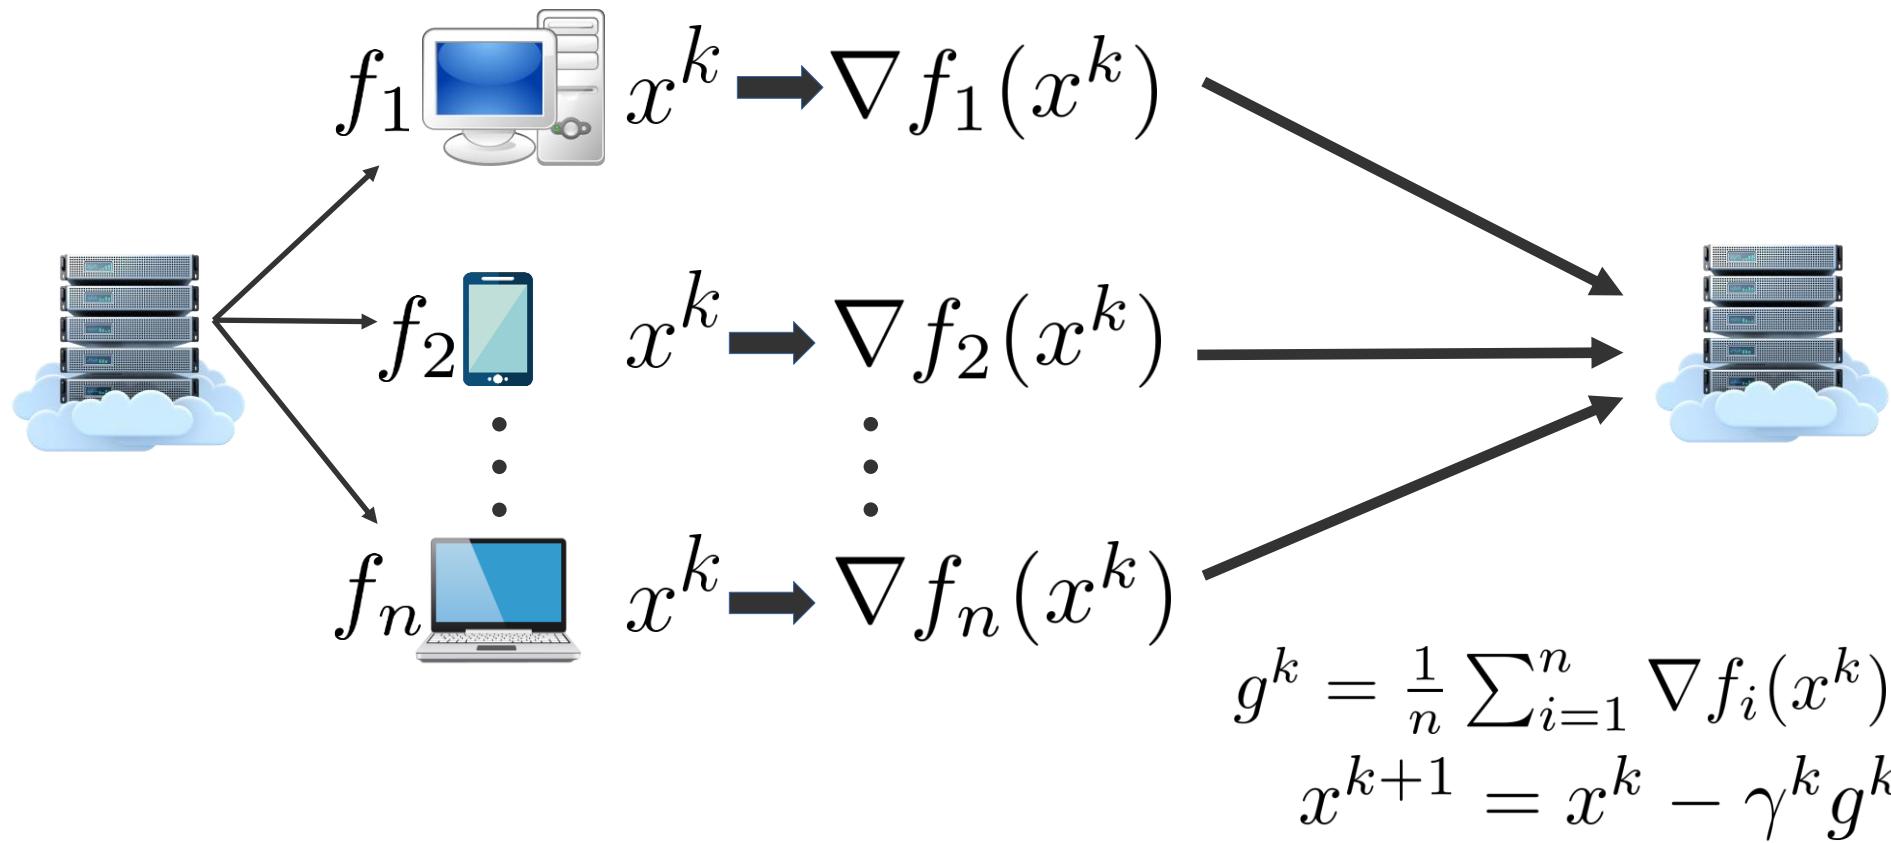
# Distributed Gradient Descent

#10



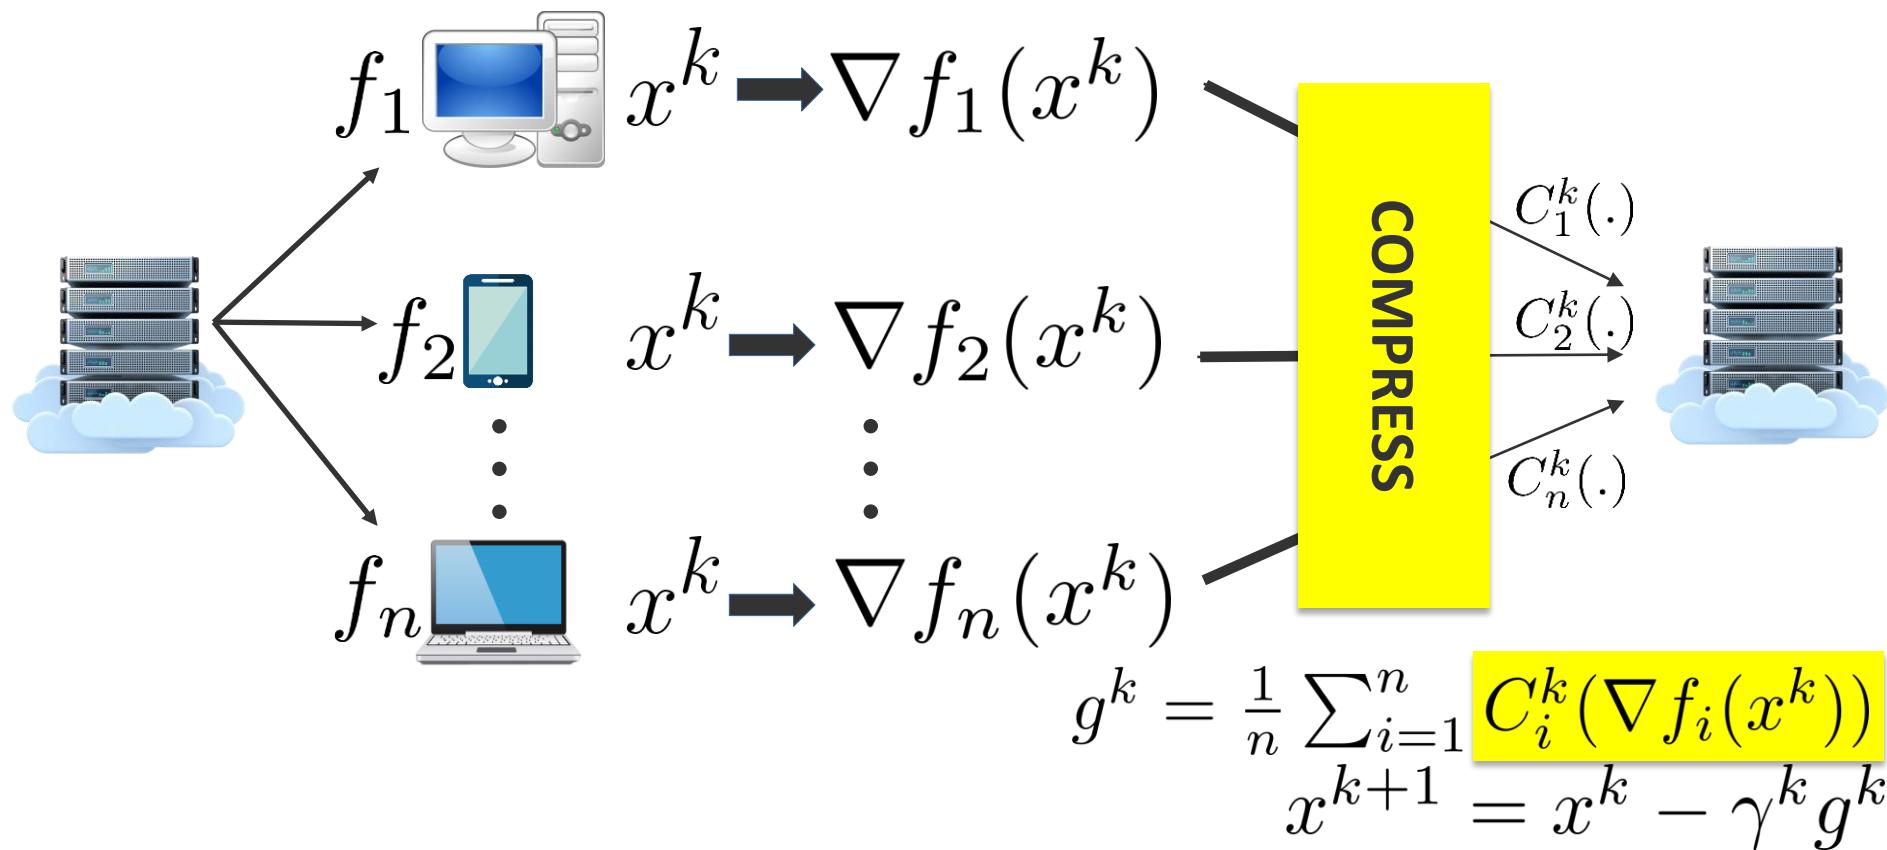
# Distributed Gradient Descent

#10



# Distributed Compressed Gradient Descent

#11



## Cost Model

Communication Complexity = (#Rounds)  $\times$  (#Bits/Round)

## Cost Model

Communication Complexity = (#Rounds) × (#Bits/Round)

## Class of Unbiased Compressors

$$\mathcal{B}^d(\omega) = \{B \mid B : \mathbb{R}^d \rightarrow \mathbb{R}^d, \quad \mathbb{E} [\|B(x) - x\|^2] \leq \omega \|x\|^2, \quad \mathbb{E} [B(x)] = x\}$$

The diagram illustrates the constraints defining the class of compressors. A red arrow points from the condition  $\omega \geq 0$  to the term  $\omega$  in the set definition, indicating that  $\omega$  must be non-negative. Another red arrow points from the condition  $\forall x \in \mathbb{R}^d$  to the term  $x$  in the set definition, indicating that the compressor must map all  $\mathbb{R}^d$  inputs to  $\mathbb{R}^d$  outputs.

## Cost Model

Communication Complexity = (#Rounds) × (#Bits/Round)

### Class of Unbiased Compressors

$$\mathcal{B}^d(\omega) = \{B \mid B : \mathbb{R}^d \rightarrow \mathbb{R}^d, \quad \mathbb{E} [\|B(x) - x\|^2] \leq \omega \|x\|^2, \quad \mathbb{E}[B(x)] = x\}$$

### Class of Contractive Compressors

$$\mathcal{C}^d(\alpha) = \{C \mid C : \mathbb{R}^d \rightarrow \mathbb{R}^d, \quad \mathbb{E} [\|C(x) - x\|^2] \leq (1 - \alpha) \|x\|^2\}$$

$$0 < \alpha \leq 1$$

$$\forall x \in \mathbb{R}^d$$

$$B \in \mathcal{B}^d(\omega) \implies C(x) := \frac{1}{\omega + 1} B(x), \quad C(x) \in \mathcal{C}^d \left( \alpha := \frac{1}{\omega + 1} \right)$$

# Compressors

## Cost Model

Communication Complexity = (#Rounds)  $\times$  (#Bits/Round)

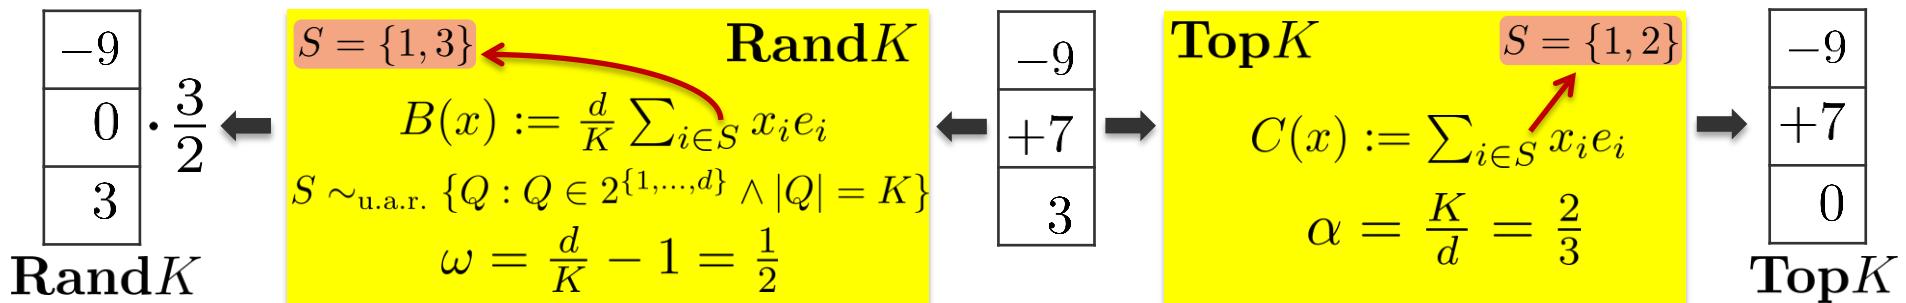
## Class of Unbiased Compressors

$$\mathcal{B}^d(\omega) = \{B \mid B : \mathbb{R}^d \rightarrow \mathbb{R}^d, \quad \mathbb{E} [\|B(x) - x\|^2] \leq \omega \|x\|^2, \quad \mathbb{E} [B(x)] = x\}$$

## Class of Contractive Compressors

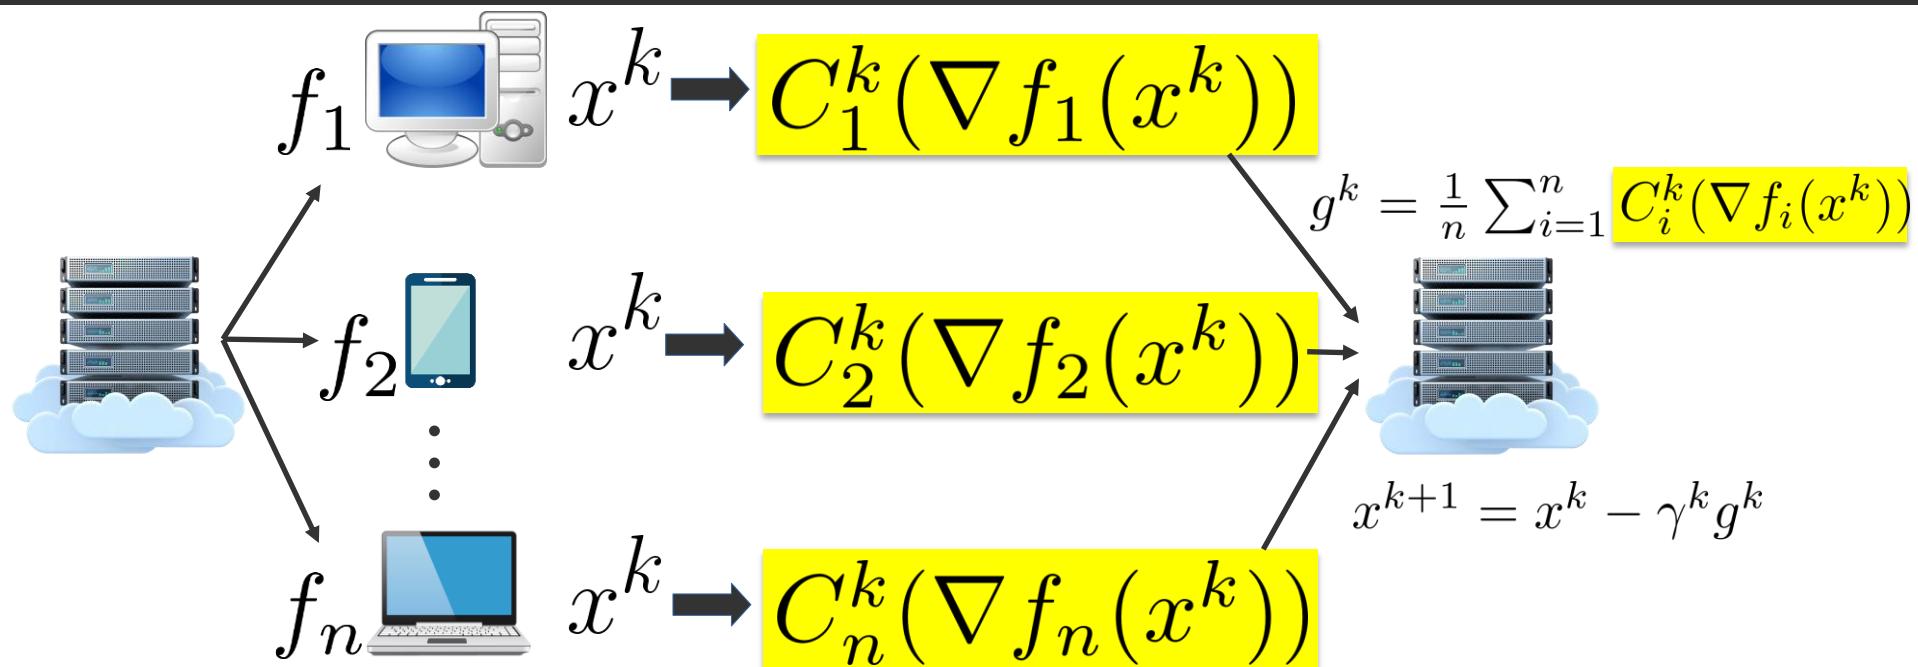
$$\mathcal{C}^d(\alpha) = \{C \mid C : \mathbb{R}^d \rightarrow \mathbb{R}^d, \quad \mathbb{E} [\|C(x) - x\|^2] \leq (1 - \alpha) \|x\|^2\}$$

## Sparsification Examples ( $d = 3, K = 2$ )



# Distributed Compressed Gradient Descent With Contractive Compressors

#13



## Distributed Compressed Gradient Descent with TopK

leads to exponential divergence even in strongly convex settings (  $n = d = 3$  )

*On Biased Compression for Distributed Learning (2023) Beznosikov et al. (Section 5.2)*

EF21 (Richtárik et al., 2021) is the theoretically fastest method that is provably correct when using contractive compressors

$$\min_{x \in \mathbb{R}^d} \left\{ f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x) \right\}$$

**Assumptions:**

1.  $f_i(x)$  are  $L_i$ -smooth, but can be non-convex
2.  $f(x)$  is  $L$ -smooth, but can be non-convex
3.  $\exists f^* > \infty$ , such that  $f(x) \geq f^*$ ,  $\forall x \in \mathbb{R}^d$

**Goal:**

Find  $\hat{x}$ :  $\mathbb{E} [\|\nabla f(\hat{x})\|^2] \leq \varepsilon^2$

EF21 (Richtárik et al., 2021) is the theoretically fastest method that is provably correct when using contractive compressors

Number of machines

$$\begin{aligned} \|\nabla f_i(x) - \nabla f_i(y)\| &\leq L_i \|x - y\| \\ \forall x, y \in \mathbb{R}^d \end{aligned}$$

$$\min_{x \in \mathbb{R}^d} \left\{ f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x) \right\}$$

Assumptions:

1.  $f_i(x)$  are  $L_i$ -smooth, but can be non-convex
2.  $f(x)$  is  $L$ -smooth, but can be non-convex
3.  $\exists f^* > \infty$ , such that  $f(x) \geq f^*$ ,  $\forall x \in \mathbb{R}^d$

Goal:

$$\text{Find } \hat{x}: \mathbb{E} [\|\nabla f(\hat{x})\|^2] \leq \varepsilon^2$$

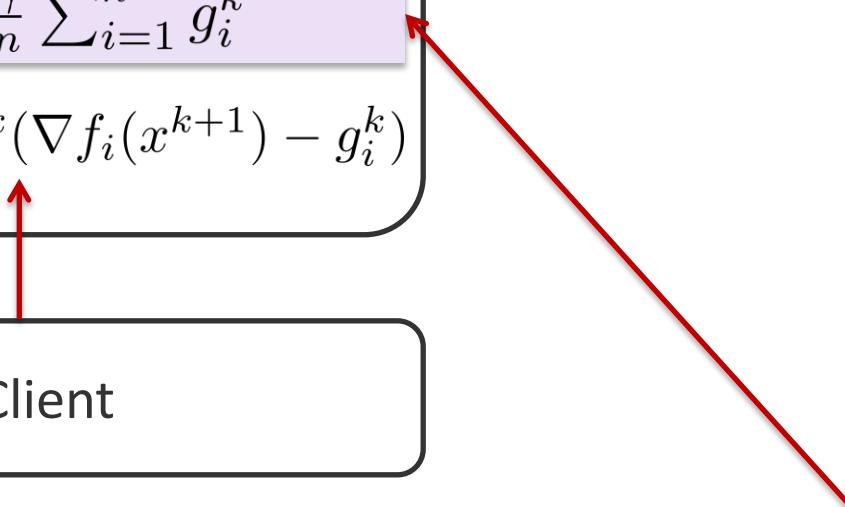
## EF21: Error Feedback 2021

$$x^{k+1} = x^k - \frac{\gamma}{n} \sum_{i=1}^n g_i^k$$

$$g_i^{k+1} = g_i^k + C_i^k (\nabla f_i(x^{k+1}) - g_i^k)$$

At Client

At Master



**EF21: Error Feedback 2021**

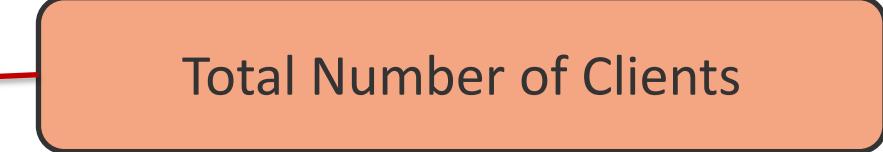
$$x^{k+1} = x^k - \frac{\gamma}{n} \sum_{i=1}^n g_i^k$$

$$g_i^{k+1} = g_i^k + C_i^k (\nabla f_i(x^{k+1}) - g_i^k)$$

Total Number of Clients

Iteration

Client



## EF21: Error Feedback 2021

$$x^{k+1} = x^k - \frac{\gamma}{n} \sum_{i=1}^n g_i^k$$

$$g_i^{k+1} = g_i^k + C_i^k (\nabla f_i(x^{k+1}) - g_i^k)$$

Communicated from  
Master to Client

Communicated from  
Client to Master

**EF21: Error Feedback 2021**

$$x^{k+1} = x^k - \frac{\gamma}{n} \sum_{i=1}^n g_i^k \quad \leftarrow$$

$$g_i^{k+1} = g_i^k + C_i^k (\nabla f_i(x^{k+1}) - g_i^k)$$

Reconstructible at the server from

- 1) received compressed messages
- 2) previous server states  $g_1^{k-1}, \dots, g_n^{k-1}$

**EF21: Error Feedback 2021**

$$x^{k+1} = x^k - \frac{\gamma}{n} \sum_{i=1}^n g_i^k$$

$$g_i^{k+1} = g_i^k + C_i^k (\nabla f_i(x^{k+1}) - g_i^k)$$

The EF21 analysis allows step size

$$0 < \gamma \leq \left( L + \sqrt{\frac{1}{n} \sum_{i=1}^n L_i^2} \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1}$$

**EF21 guarantees**

$$\approx \frac{1}{\alpha}, \alpha \in (0, 0.5)$$

$$\mathbb{E} [\|\nabla f(\hat{x}^T)\|^2] \leq \frac{2(f(x^0) - f^\star)}{\gamma T} + \frac{G^0}{\theta(\alpha)T}$$

$$\theta(\alpha) := 1 - \sqrt{1 - \alpha}, \beta(\alpha) := \frac{1 - \alpha}{1 - \sqrt{1 - \alpha}}$$

$$\alpha = \frac{K}{d} \text{ for Top } K \text{ compressor}$$

## EF21: Error Feedback 2021

$$x^{k+1} = x^k - \frac{\gamma}{n} \sum_{i=1}^n g_i^k$$

$$g_i^{k+1} = g_i^k + C_i^k (\nabla f_i(x^{k+1}) - g_i^k)$$

The EF21 analysis allows step size

$$0 < \gamma \leq \left( L + \sqrt{\frac{1}{n} \sum_{i=1}^n L_i^2} \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1}$$

$$f^\star = \inf_{x \in \mathbb{R}^d} f(x)$$

$$G^0 := \sum_{i=1}^n \frac{1}{n} \|g_i^0 - \nabla f_i(x^0)\|^2$$

$$\mathbb{E} [\|\nabla f(\hat{x}^T)\|^2] \leq \frac{2(f(x^0) - f^\star)}{\gamma T} + \frac{G^0}{\theta(\alpha)T}$$

$$\hat{x}^T \sim_{\text{u.a.r.}} \{x^0, \dots, x^{T-1}\}$$

**Total Iterations**

$$\theta(\alpha) := 1 - \sqrt{1 - \alpha}, \beta(\alpha) := \frac{1 - \alpha}{1 - \sqrt{1 - \alpha}}$$

$$\alpha = \frac{K}{d} \text{ for Top } K \text{ compressor}$$

## The best step size for EF21

$$\gamma = \left( L + \sqrt{\frac{1}{n} \sum_{i=1}^n L_i^2} \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1}$$

## The best step size for EF21

$$\gamma = \left( L + \sqrt{\frac{1}{n} \sum_{i=1}^n L_i^2} \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1}$$

Can we decrease it?

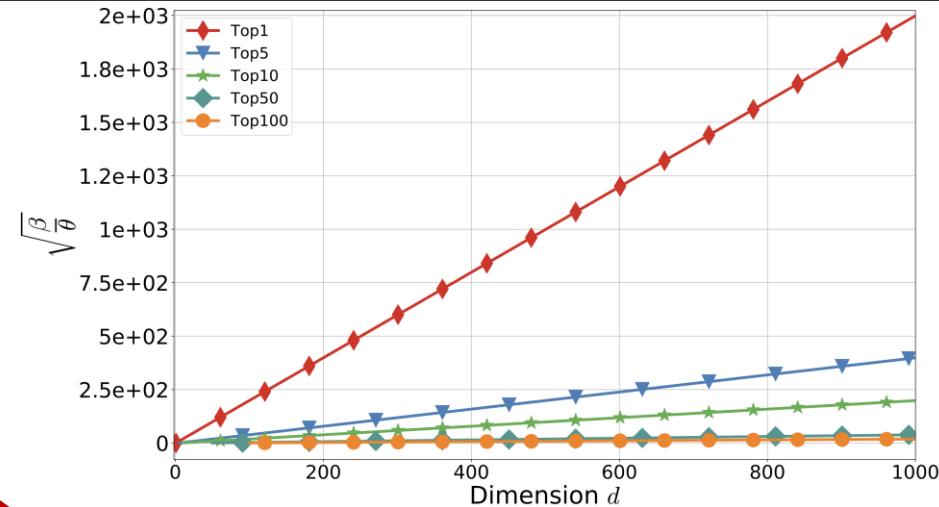
This is already very important

$$L \leq \underbrace{\frac{1}{n} \sum L_i}_{\text{AM}} \leq \underbrace{\sqrt{\frac{1}{n} \sum L_i^2}}_{\text{QM}}$$

## The best step size for EF21

$$\gamma = \left( L + \sqrt{\frac{1}{n} \sum_{i=1}^n L_i^2} \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1}$$

Can we decrease it?



And it can be arbitrarily big for  
Top $K$  with  $K \ll d$

The best step size for EF21

$$\gamma = \left( L + \sqrt{\frac{1}{n} \sum_{i=1}^n L_i^2} \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1}$$

We improved the step size in 3 different ways to

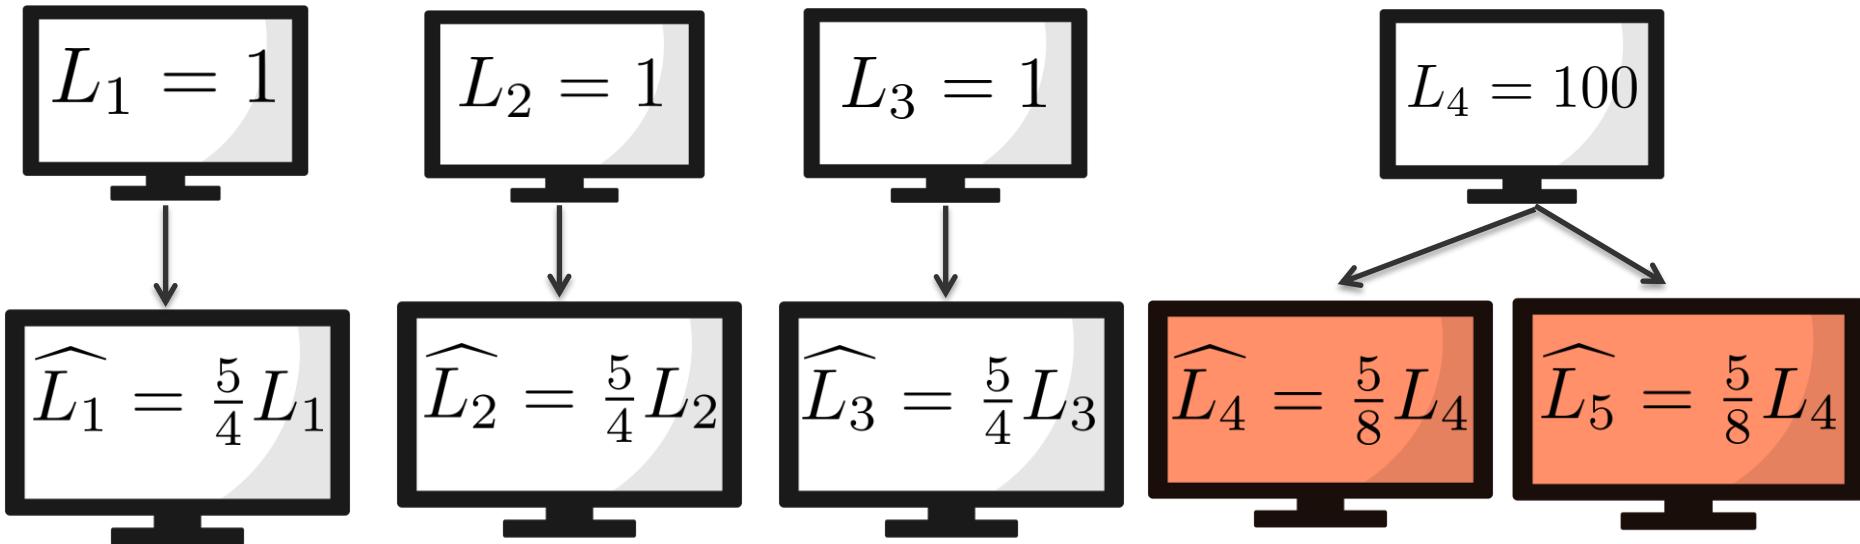
$$\gamma = \left( L + \frac{1}{n} \sum_{i=1}^n L_i \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1}$$

# Example-Driven Reformulation

#18

$$L_{AM} := \frac{(1+1+1+100)}{4} = 25.75$$

$$L_{QM} = \sqrt{\frac{(1+1+1+100 \cdot 100)}{4}} = \sqrt{2500.75}$$



$$\hat{L}_{AM} = \frac{3 \cdot (5/4) + 2 \cdot (500/8)}{5} = 25.75$$

$$\hat{L}_{QM} = \sqrt{\frac{3 \cdot (5/4)^2 + 2 \cdot (500/8)^2}{5}} = \sqrt{1563}$$

QM changed, even AM is the same !

# Ch3: EF21 Reloaded (Approach 1)

#19  
• ○ ○



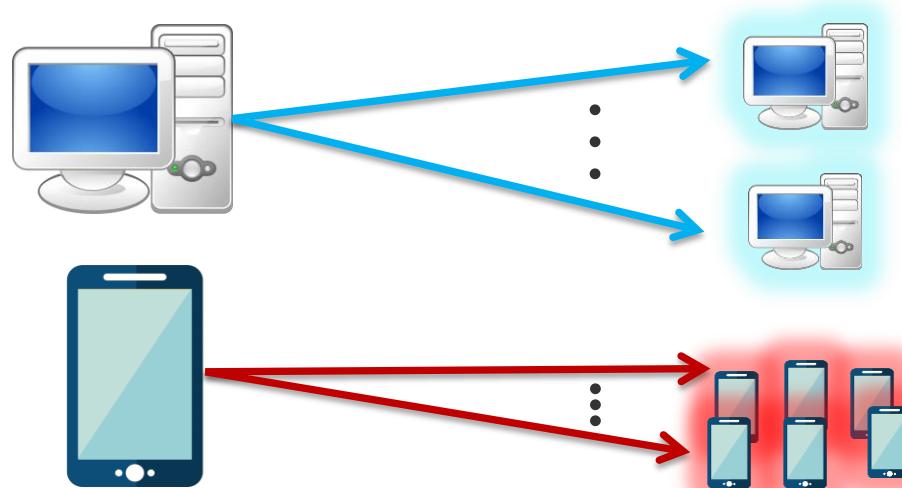
$n$  clients with  $f_i(x)$

$$f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x)$$

VIRTUAL CLONING OF EF21 CLIENTS

# Ch3: EF21 Reloaded (Approach 1)

#19  
• • •

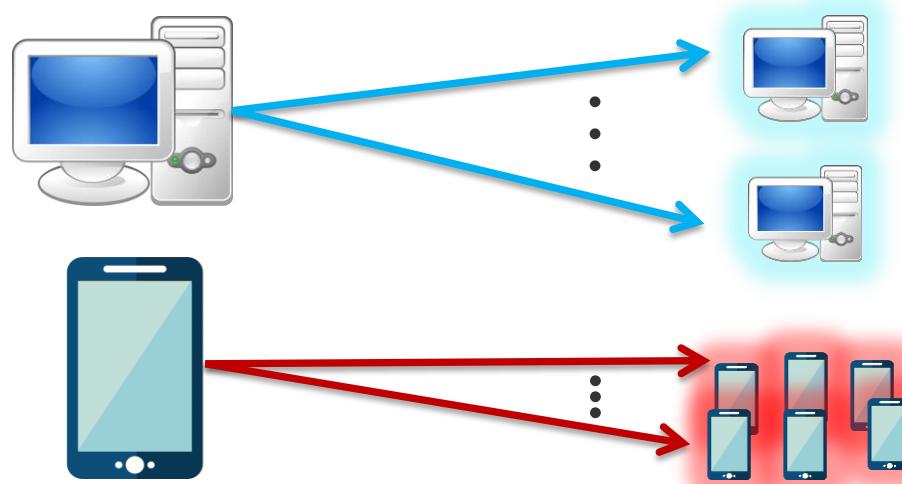


$n$  clients with  $f_i(x)$  Client  $i$  cloned  $N_i$  times

$$f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x) \quad N := \sum_{i=1}^n N_i$$

**VIRTUAL CLONING OF EF21 CLIENTS**

# Ch3: EF21 Reloaded (Approach 1)

#19  
..

$n$  clients with  $f_i(x)$  Client  $i$  cloned  $N_i$  times

$$f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x)$$

$$N := \sum_{i=1}^n N_i$$



$$\widehat{f}_{ij}(x) = \frac{N}{nN_i} f_i(x)$$

$$\implies L_{ij} = \frac{N}{nN_i} L_i$$

$$\widehat{f}(x) := \frac{1}{N} \sum_{i=1}^n \sum_{j=1}^{N_i} \widehat{f}_{ij}(x)$$

**VIRTUAL CLONING OF EF21 CLIENTS**

# Ch3: EF21 Reloaded (Approach 1)

#20  
• °

$$\widehat{f}(x) := \frac{1}{N} \sum_{i=1}^n \sum_{j=1}^{N_i} \widehat{f}_{ij}(x)$$



$$\widehat{f}_{ij}(x) = \frac{N}{nN_i} f_i(x), L_{ij} = \frac{N}{nN_i} L_i$$

$$N := \sum_{i=1}^n N_i$$

Each client  $i$  cloned  $N_i$  times



**VIRTUAL CLONING OF EF21 CLIENTS**

# Ch3: EF21 Reloaded (Approach 1)

#20  
• •

$$\widehat{f}(x) := \frac{1}{N} \sum_{i=1}^n \sum_{j=1}^{N_i} \widehat{f}_{ij}(x)$$

$$L_{ij} = \frac{N}{n N_i} L_i$$

$$M(N_1, \dots, N_n) := \sqrt{\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^{N_i} L_{ij}^2} = \frac{1}{n} \sqrt{\sum_{i=1}^n \frac{L_i^2}{N_i/N}}$$

$$\min_{N_i \in \mathbb{R}, N_i > 0, \sum_{i=1}^n N_i/N = 1} M(N_1, \dots, N_n) = \frac{\sum_{i=1}^n L_i}{n}$$

$$\frac{\sum_{i=1}^n L_i}{n} \leq M(\lceil L_1/L_{\text{AM}} \rceil, \dots, \lceil L_n/L_{\text{AM}} \rceil)) \leq (\frac{1}{n} \sum_{i=1}^n L_i) \sqrt{2}$$

$$N_1^*, \dots, N_n^*$$



**Good:** We reduced QM to AM (up to the factor  $\sqrt{2}$ )

**Bad:** We need to increase number of workers  $n \rightarrow N, n \leq N \leq 2n$

$$\gamma \approx \left( L + \frac{1}{n} \sum_{i=1}^n L_i \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1} \quad N \geq n$$

**Good:** We reduced QM to AM (up to the factor  $\sqrt{2}$ )

**Bad:** We need to increase number of workers  $n \rightarrow N, n \leq N \leq 2n$

$$\gamma \approx \left( L + \frac{1}{n} \sum_{i=1}^n L_i \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1} \quad N \geq n$$

## Assumptions:

B1. Initial shifts for all clones are identical

B2. The compressors are deterministic

⇒ Under these assumptions, the cloning mechanism can be reformulated as a new **EF21-W**

# Ch3: EF21 Reloaded (Approach 2)

**Good:** We reduced QM to AM (up to the factor  $\sqrt{2}$ )

**Bad:** We need to increase number of workers  $n \rightarrow N, n \leq N \leq 2n$

$$\gamma \approx \left( L + \frac{1}{n} \sum_{i=1}^n L_i \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1} \quad N \geq n$$

**Assumptions:**

B1. Initial shifts for all clones are identical

B2. The compressors are deterministic

⇒ Under these assumptions, the cloning mechanism can be reformulated as a new **EF21-W**

$$x^{k+1} = x^k - \frac{\gamma}{n} \sum_{i=1}^n g_i^k$$

$$g_i^{k+1} = g_i^k + C_i^k (\nabla f_i(x^{k+1}) - g_i^k)$$



$$x^{k+1} = x^k - \frac{\gamma}{N} \sum_{i=1}^n \sum_{j=1}^{N_i} g_{ij}^k$$

$$g_{ij}^{k+1} = g_{ij}^k + C_i^k (\nabla f_{ij}(x^{k+1}) - g_{ij}^k)$$



$$x^{k+1} = x^k - \gamma \sum_{i=1}^n w_i g_i^k$$

$$g_i^{k+1} = g_i^k + C_i^k \left( \frac{1}{n w_i} \nabla f_i(x^{k+1}) - g_i^k \right)$$

$$w_i := \frac{L_i}{\frac{1}{n} \sum_{i=1}^n L_i}$$

EF21-W

# Ch3: EF21 Reloaded (Approach 2)

**Good:** We reduced QM to AM (up to the factor  $\sqrt{2}$ )

**Bad:** We need to increase number of workers  $n \rightarrow N, n \leq N \leq 2n$

$$\gamma \approx \left( L + \frac{1}{n} \sum_{i=1}^n L_i \times \sqrt{\frac{\beta(\alpha)}{\theta(\alpha)}} \right)^{-1} \quad N \geq n$$

## Assumptions:

B1. Initial shifts for all clones are identical

B2. The compressors are deterministic

⇒ Under these assumptions, the cloning mechanism can be reformulated as a new **EF21-W**

$$\begin{aligned} x^{k+1} &= x^k - \gamma \sum_{i=1}^n w_i g_i^k \\ g_i^{k+1} &= g_i^k + C_i^k \left( \frac{1}{n w_i} \nabla f_i(x^{k+1}) - g_i^k \right) \end{aligned} \quad w_i := \frac{L_i}{\frac{1}{n} \sum_{i=1}^n L_i}$$

Our analysis reveals that assumptions (B1) and (B2) are not required

The analysis of **EF21-W** reveals that the original EF21 analysis requires modification for the quantity  $G^t$

$$G_i^t := \|g_i^t - \nabla f_i(x^t)\|^2 \quad G^t := \sum_{i=1}^n \frac{1}{n} G_i^t$$

$$\begin{aligned} G_i^t &:= \|g_i^t - \frac{\nabla f_i(x^t)}{nw_i}\|^2 \quad G^t := \sum_{i=1}^n w_i G_i^t \\ w_i &:= \frac{L_i}{\frac{1}{n} \sum_{i=1}^n L_i} \end{aligned}$$

It motivated us to analyze the original EF21 and discover:

**Incorporating weights into the original EF21 analysis improves the rate !!**

# Federated Learning Challenges Addressed in the Thesis #23

## Theoretical Work

1. Data Heterogeneity

2. Device Heterogeneity

3. Communication Bottleneck

4. Privacy

5. Software

## Theory-Inspired Practical Work

## Practical Work

### Ch1: Introduction



**Ch3: EF21-W**  
Richtárik et al., 2024

**Ch3: EF21-W**  
Richtárik et al., 2024

**Ch4: DCGD/PERMK/AES**  
Burlachenko et al., 2023

**Ch5: PAGE Extensions**  
Tyurin et al., 2023

**Ch6:**  
**Compressed L2GD**  
Bergou et al., 2023

**Ch7: Unlocking FedNL**  
Burlachenko and Richtárik, 2024

**Ch8: BurTorch**  
Burlachenko & Richtárik, 2025

**Ch2: FL\_PyTorch**  
Burlachenko et al., 2021

**Ch9: Concluding Remarks: Summary and Future Research**

# Main Tools for Privacy Guarantees in FL

#24

## Trusted Execution Environments (TEE)

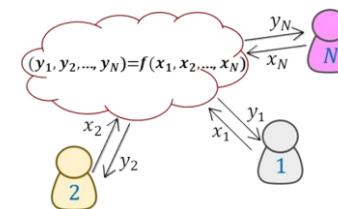
Protects the execution environment from illegal intervention



## Differential Privacy (DP)

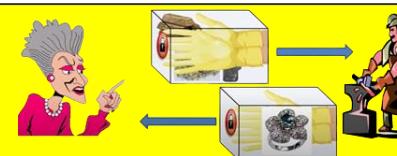


Protects output of algorithm so that users' data are not leaking after execution



## Secure Multi-Party Computation (MPC)

Protects inputs of algorithm at the cost of communication



## Homomorphic Encryption (HE)

Computation on encrypted data without revealing inputs or outputs

# Homomorphic Encryption (HE)

#25

Homomorphism of two groups  $G_1$  and  $G_2$  is a mapping  $f : G_1 \rightarrow G_2$

$$f(x * y) = f(x) * f(y), \quad \forall x, y \in G_1$$



## Homomorphic Encryption:

Computation on encrypted data without revealing inputs or outputs



## Homomorphic Encryption In Action:

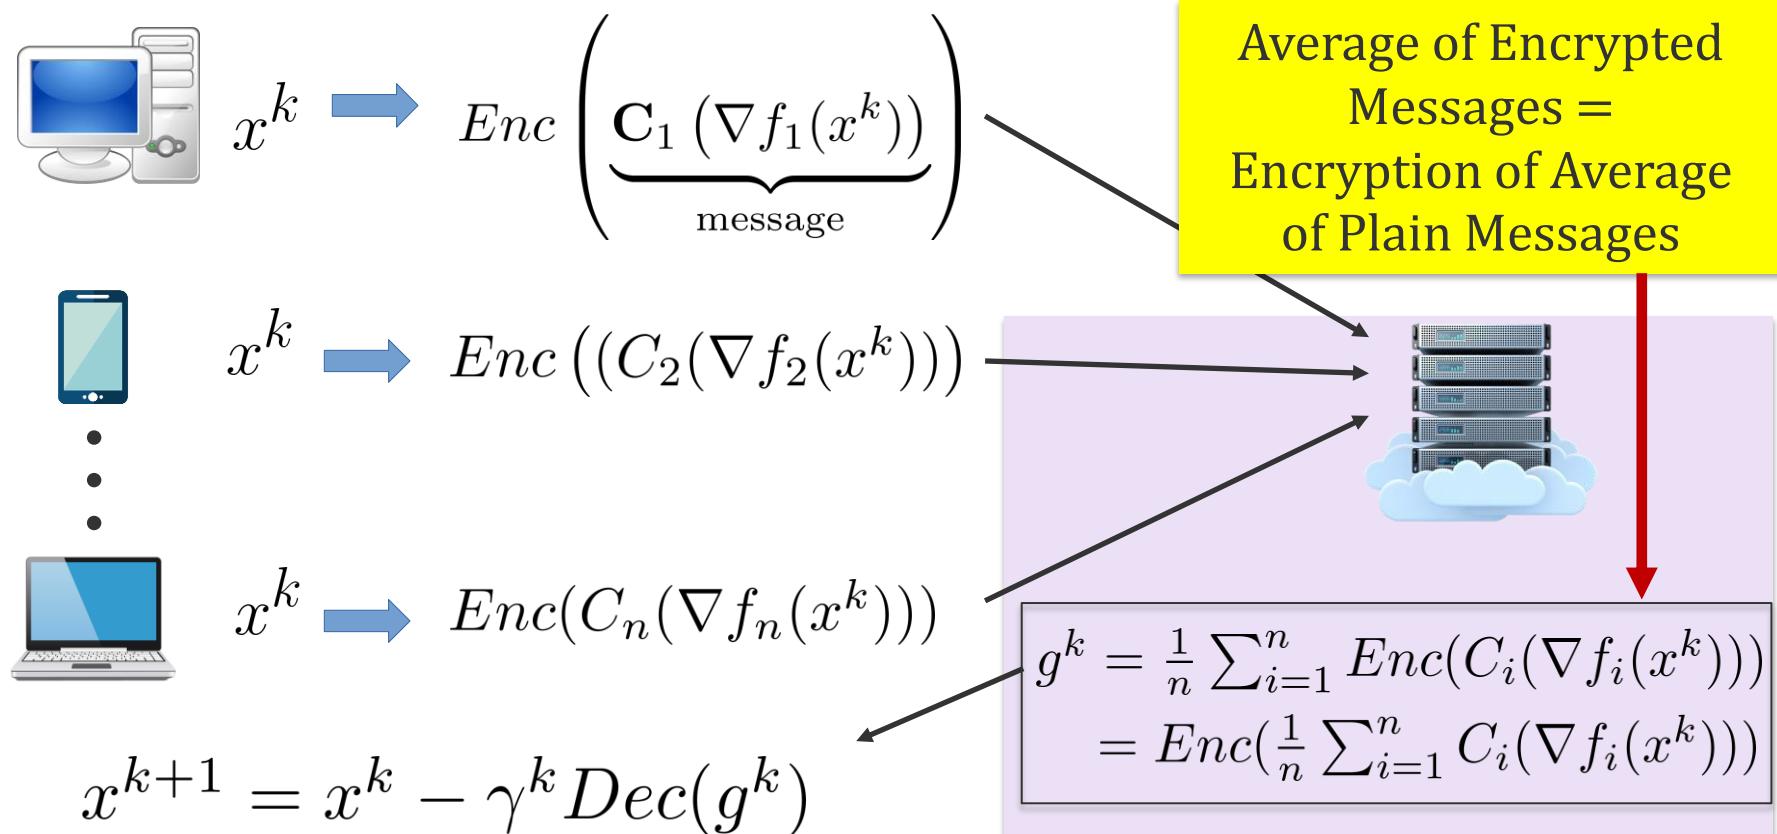
1. Any device with the **public key** can perform computations on encrypted data
2. Only the holder of the **private key** can decrypt the result

## Cheon-Kim-Kim-Song (CKKS, 2017):

- a. The **CKKS** scheme supports approximate arithmetics on real and complex dense vectors and is considered as SOTA in this class
- b. **CKKS (and HE in general)** is more complex primitive than classical block ciphers (e.g. AES-based), relying on entirely different mathematical foundations

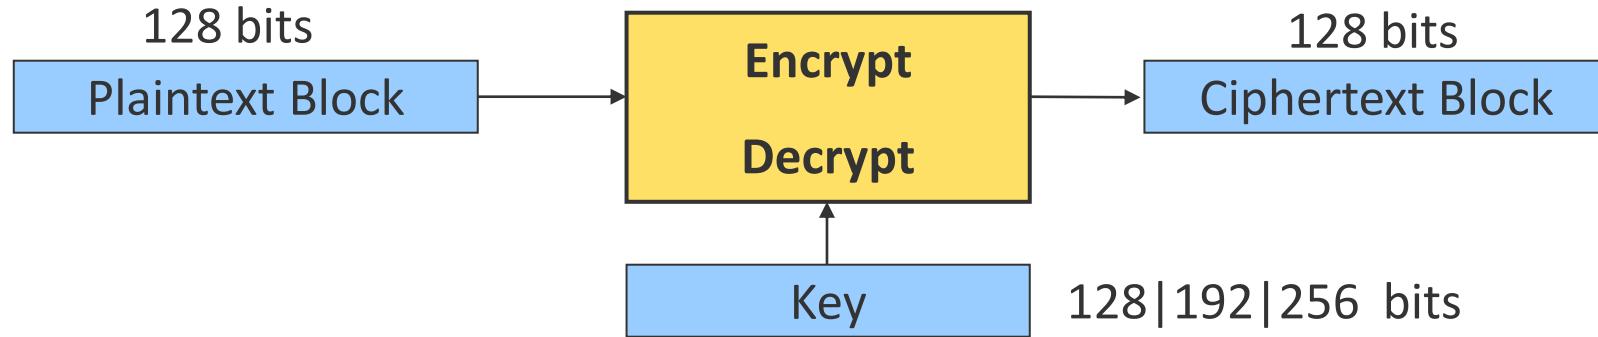
# Distributed Compressed Gradient Descent with Homomorphic Encryption (HE)

#26



# Classical Cryptography: AES Block Cipher

#27



## AES (2001) Block Cipher

- Maps deterministically and with reversible operations input (128 bits) into output (128 bits)
- Has hardware support (Intel Westmere, AMD Bulldozer, ARM Cortex-A53)
- AES is a strong cryptographic primitive, widely trusted as a secure **PRP**

### A secure pseudorandom permutation (PRP)

produces permutations that are computationally indistinguishable from uniformly random permutations by any *known* polynomial-time algorithm



DP, HE, MPC, TEE...

# But where is Classical Cryptography?

#28

Researchers from 2020 – 2024 consistently argue that applying symmetric-key encryption like AES or DES in FL is **unsuitable, challenging, not feasible**



**Secure, Privacy-Preserving, and Federated Machine Learning in Medical Imaging,**  
G. Kaassis et al. (2020) Nature Machine Intelligence

**Private Artificial Intelligence: Machine Learning on Encrypted Data,**  
Kristin E. Lauter (2022) SIAM

**Cybersecurity English Accelerated Encrypted Execution of General Purpose Applications,**  
V. Joseph et al (2023) NVIDIA Blog

**FedSHE: Privacy-Preserving and Efficient Federated Learning with Adaptive Segmented CKKS Homomorphic Encryption,**  
Pan Y. et al. (2024) Cybersecurity

**Revisiting Fully Homomorphic Encryption Schemes for Privacy-Preserving Computing,**  
N. Jain et al. (2024) Emerging Technologies and Security in Cloud Computing

# Ch4: DCGD/PermK/AES (2023)

#29



Home > Conferences > CONEXT > Proceedings > DistributedML '23 > Federated Learning is Better with Non-Homomorphic Encryption

RESEARCH ARTICLE OPEN ACCESS

X in

## Federated Learning is Better with Non-Homomorphic Encryption

Authors: Konstantin Burdachenko, Abdulmajeed Alrowithi, Fahad Ali Albalawi, Peter Rischárik [Authors Info & Claims](#)

DistributedML '23: Proceedings of the 4th International Workshop on Distributed Machine Learning • December 2023 • Pages 49–84 • <https://doi.org/10.1145/9630048.3630182>

Published: 05 December 2023 [Publication History](#)

99 0 202

**ABSTRACT**

DistributedML '23: Proceedings of the 4th International Workshop on Distributed Machine Learning • December 2023 • Pages 49–84 • <https://doi.org/10.1145/9630048.3630182>

Traditional AI methodologies necessitate centralized data collection, which becomes impractical when facing problems with network communication, data privacy, or storage capacity. Federated Learning (FL) offers a paradigm that empowers distributed AI model training without collecting raw data. There are different choices for providing privacy during FL training. One of the popular methodologies is employing Homomorphic Encryption (HE) – a breakthrough in privacy-preserving computation from Cryptography. However, these methods have a price in the form of extra computation and memory footprint. To resolve these issues, we propose an innovative framework that synergizes permutation-based compressors with Classical Cryptography, even though employing Classical Cryptography was assumed to be impossible in the past in the context of FL. Our framework offers a way to replace HE with cheaper Classical Cryptography primitives which provides security for the training process. It fosters asynchronous communication and provides flexible deployment options in various communication topologies.

← Previous Next →

REFERENCES

Index Terms

Recommendations

Comments

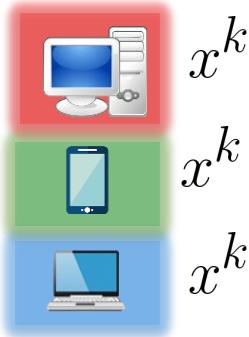
ACM DIGITAL LIBRARY



Distributed  
Compressed  
Gradient Descent

Permutated  
Correlated  
Compressors

Advanced Encryption  
Standard

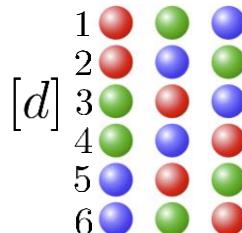


## PermK Compressors (Rafał Szlendak, et al. 2021)



$$x^k \rightarrow g_i^k = B_i(\nabla f_i(x^k))$$

$$x^k \rightarrow$$



$$x^k \rightarrow$$

$$k = 1, 2, 3, \dots$$



Green user uses coordinates 3, 4 from  $[d] = \{1, \dots, 6\}$

## Example



$$x^k \rightarrow g_i^k = B_i(\nabla f_i(x^k))$$

$$x^k \rightarrow [d]$$

1	2	3	4	5	6
2	3	4	5	6	1
3	4	5	6	1	2
4	5	6	1	2	3
5	6	1	2	3	4
6	1	2	3	4	5

$$x^k \rightarrow k = 1, 2, 3, \dots$$

$$\nabla f_{1,2,3}(x^k)$$

0.1	1.1	2.1
0.2	1.2	2.2
0.3	1.3	2.3
0.4	1.4	2.4
-0.5	1.5	2.5
0.6	1.6	2.6



$$3 \times$$

$$g_1^k$$

0.1	0	0
0.2	0	0
0	1.3	0
0	1.4	0
0	0	2.5
0	0	2.6

$$g_2^k$$

0	0	0
0	0	0
1.3	0	0
1.4	0	0
0	2.5	0
0	0	2.6

$$g_3^k$$

0	0	0
0	0	0
0	1.3	0
0	1.4	0
0	0	2.5
0	0	2.6

Training Iteration

Scaling is needed to preserve unbiasedness



$$\begin{aligned}
 x^k &\rightarrow g_i^k = B_i(\nabla f_i(x^k)) \rightarrow g_1^k \\
 x^k &\rightarrow [d] \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ 6 \end{matrix} \begin{matrix} \text{red} \\ \text{red} \\ \text{green} \\ \text{green} \\ \text{blue} \\ \text{blue} \end{matrix} \rightarrow g_2^k \\
 x^k &\rightarrow [d] \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ 6 \end{matrix} \begin{matrix} \text{green} \\ \text{green} \\ \text{red} \\ \text{red} \\ \text{blue} \\ \text{blue} \end{matrix} \rightarrow g_n^k
 \end{aligned}$$

$k = 1, 2, 3, \dots$

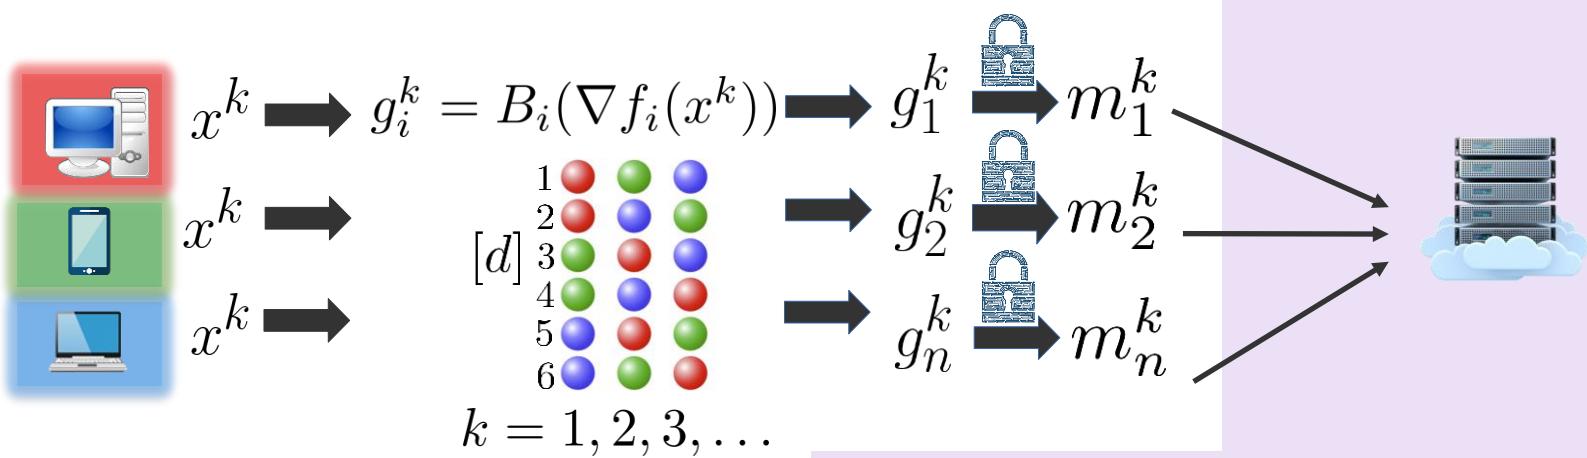
### Algebraic Properties of PermK

$$\begin{aligned}
 \mathbb{E} \left[ \frac{1}{n} \sum_{i=1}^n B_i(v_i) \right] &= \frac{1}{n} \sum_{i=1}^n v_i & \forall v_i \in \mathbb{R}^d \\
 \mathbb{E} \left[ \left\| \frac{1}{n} \sum_{i=1}^n B_i(v_i) - \frac{1}{n} \sum_{i=1}^n v_i \right\|^2 \right] &\leq \frac{1}{n} \sum_{i=1}^n \|v_i\|^2 - \left\| \frac{1}{n} \sum_{i=1}^n v_i \right\|^2
 \end{aligned}$$

# DCGD/PermK/AES

#30

○ ○ ○ ○ ● ○ ○

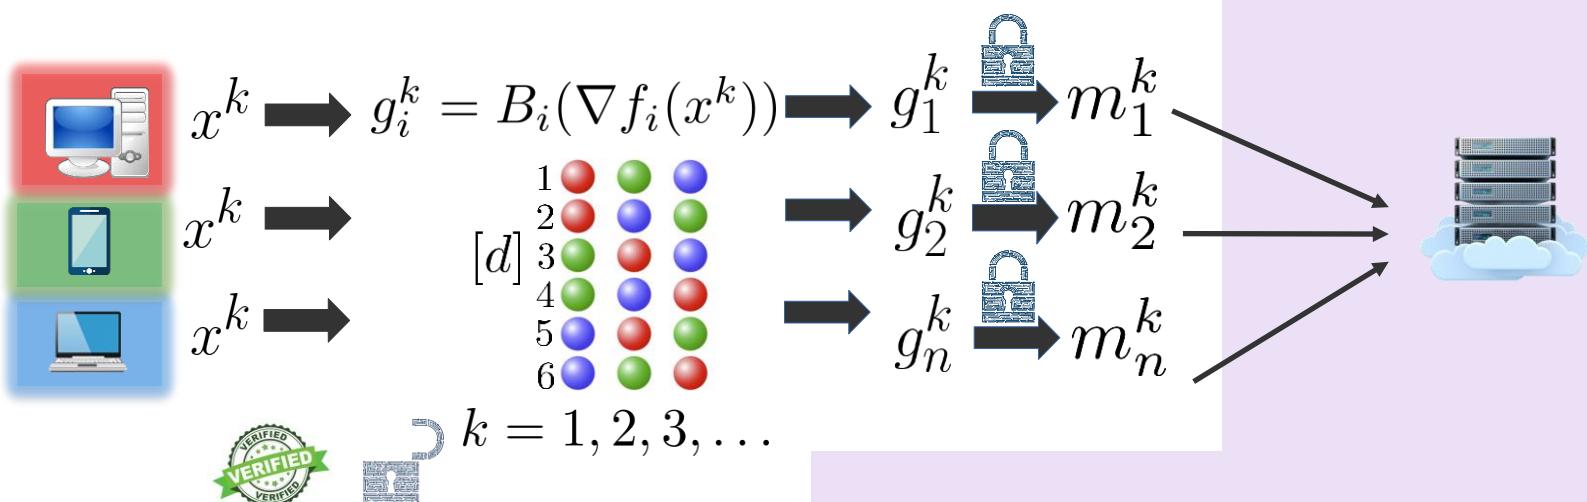


$$g^k = \sum_{i=1}^n m_i^k = \text{concat}(m_1^k, \dots, m_n^k)$$

# DCGD/PermK/AES

#30

.....



$$x^{k+1} = x^k - \gamma^k \cdot \frac{1}{n} [Dec(m_1^k), \dots, Dec(m_n^k)]^\top \leftarrow g^k = \sum_{i=1}^n m_i^k = \text{concat}(m_1^k, \dots, m_n^k)$$

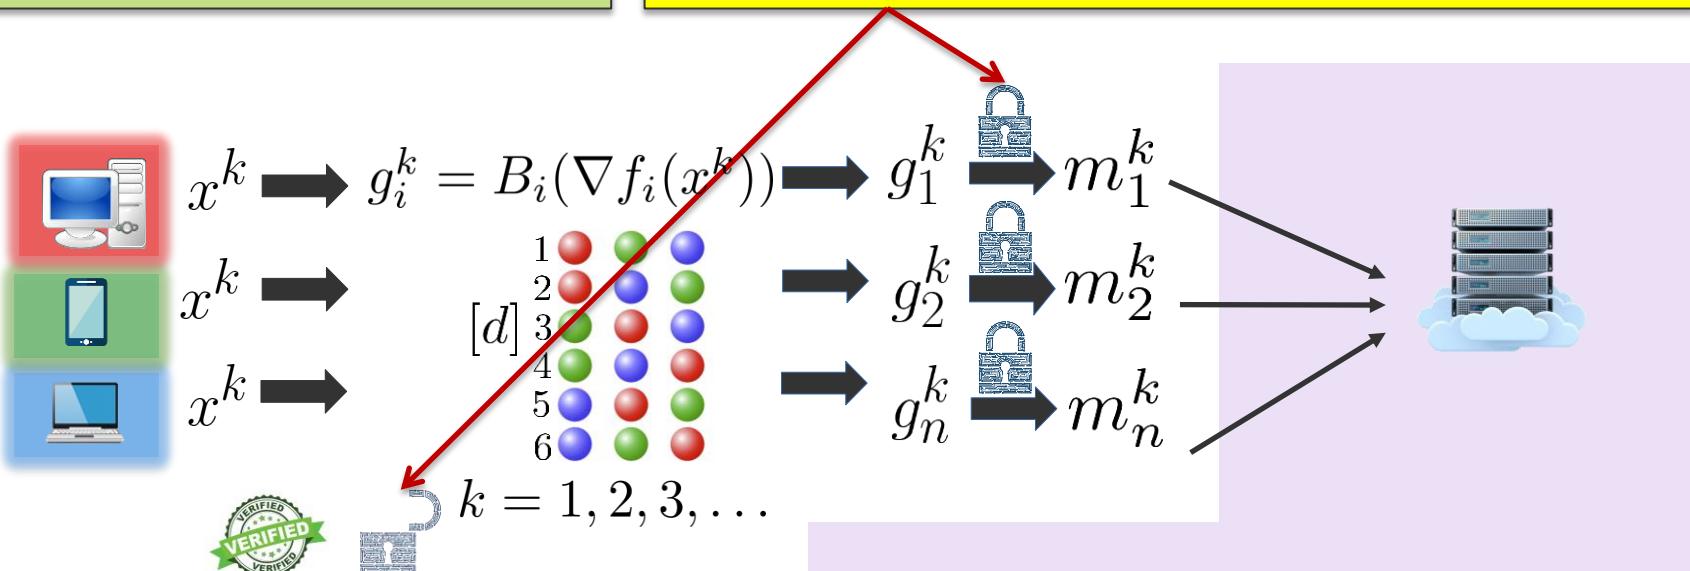
# DCGD/PermK/AES

#30

.....

HE/CKKS for AES-128 security level  
requires key(s) with sizes 420 000 bytes

For AES-128 the key size is 128 bits (16 bytes)



$$x^{k+1} = x^k - \gamma^k \cdot \frac{1}{n} [Dec(m_1^k), \dots, Dec(m_n^k)]^\top \leftarrow g^k = \sum_{i=1}^n m_i^k = \text{concat}(m_1^k, \dots, m_n^k)$$

# DCGD/PermK/AES vs HE

$$\text{For } g_i^k = B_i(\nabla f_i(x^k))$$

$$\text{Enc}\left(\frac{1}{n} \sum_{i=1}^n g_i^k\right) = \frac{1}{n} \text{Concat}(\text{Enc}(g_1^k), \text{Enc}(g_2^k), \dots, \text{Enc}(g_n^k))$$

NEW

- Only compatible with specific
- + Does not introduce numerical errors
- + Low memory overhead from AES

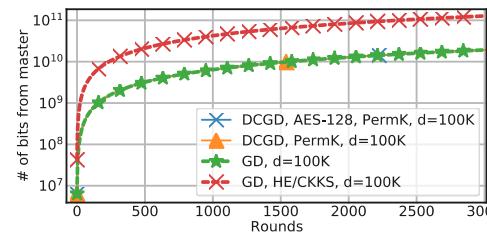
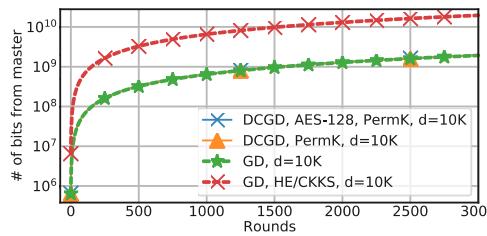
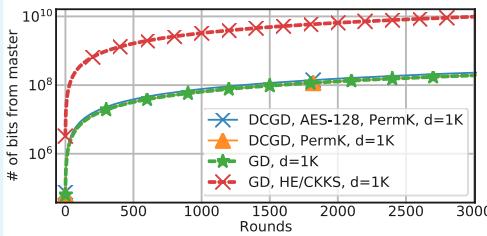
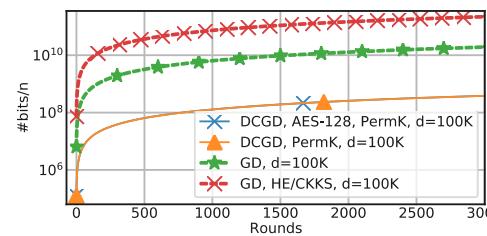
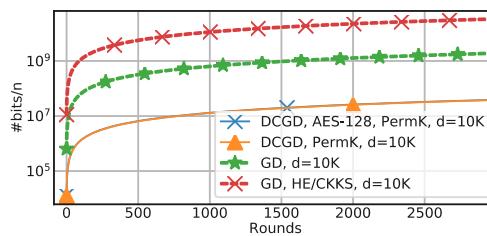
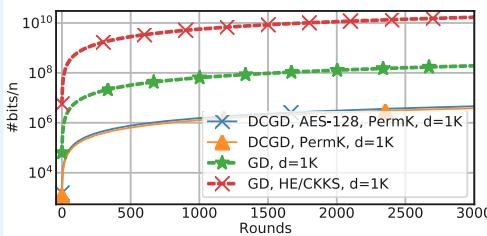
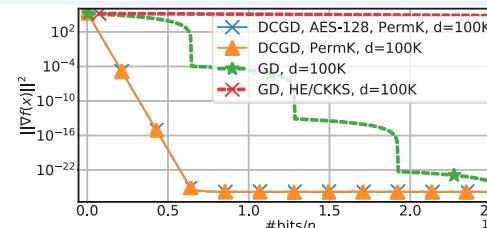
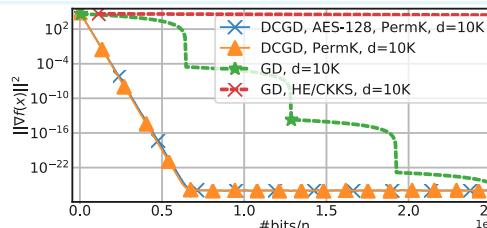
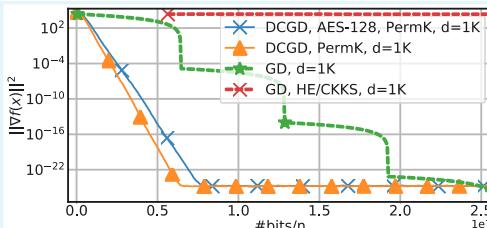
$$\forall g_i^k \in \mathbb{R}^d$$

$$\text{Enc}\left(\frac{1}{n} \sum_{i=1}^n g_i^k\right) = \frac{1}{n} \sum_{i=1}^n \text{Enc}(g_i^k)$$

CLASSICAL HE

- + Works with arbitrarily
- Introduces numerical errors
- High memory overhead

# DCGD/Permk/AES vs DCGD/PermK vs GD vs GD/CKKS #32



(a)  $d = 1000$

(b)  $d = 10000$

(c)  $d = 100000$

Linear regression

Interpolation regime

$n = 50, n_i = 12$

Compute FP64

Tuned  $\gamma = 0.007$  for DCGD

Theoretical  $\gamma$  for GD

$$\min_{x \in \mathbb{R}^d} \left\{ f(x) := \frac{1}{4} \sum_{i=1}^4 f_i(x) \right\}$$

## Gradient Descent

$$g^k = \sum_{i=1}^4 \nabla f_i(x)$$

$$x^{k+1} = x^k - \gamma \cdot \frac{1}{n} \cdot g^k$$

Requires synchronization among clients

## DCGD/PermK/AES

$$g^k = [\textcolor{red}{p_1}, p_2, p_3, p_4]^\top$$

$$x_{\text{parts}:2,3,4}^{k+1} = x_{\text{parts}:2,3,4}^k - \gamma \cdot \frac{1}{n} \cdot [g^k]_{\text{parts}:2,3,4}$$

Start forward pass for next iteration with partial  $x^{k+1}$

Wait for straggler (client number  $\#1$ )

# Federated Learning Challenges Addressed in the Thesis #34

## Theoretical Work

1. Data Heterogeneity

2. Device Heterogeneity

3. Communication Bottleneck

4. Privacy

5. Software

## Theory-Inspired Practical Work

## Practical Work

### Ch1: Introduction

Ch3: EF21-W  
Richtárik et al., 2024



Ch3: EF21-W  
Richtárik et al., 2024

Ch4: DCGD/PERMK/AES  
Burlachenko et al., 2023

The  $f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x)$  holds a surprising property

Ch2: FL\_PyTorch

Burlachenko et al., 2021

Ch5: PAGE Extensions  
Tyurin et al., 2023

Ch6:  
Compressed L2GD  
Bergou et al., 2023

Ch7: Unlocking FedNL  
Burlachenko and Richtárik, 2024

Ch7: Unlocking FedNL  
Burlachenko & Richtárik, 2024

Ch8: BurTorch

Burlachenko & Richtárik, 2025

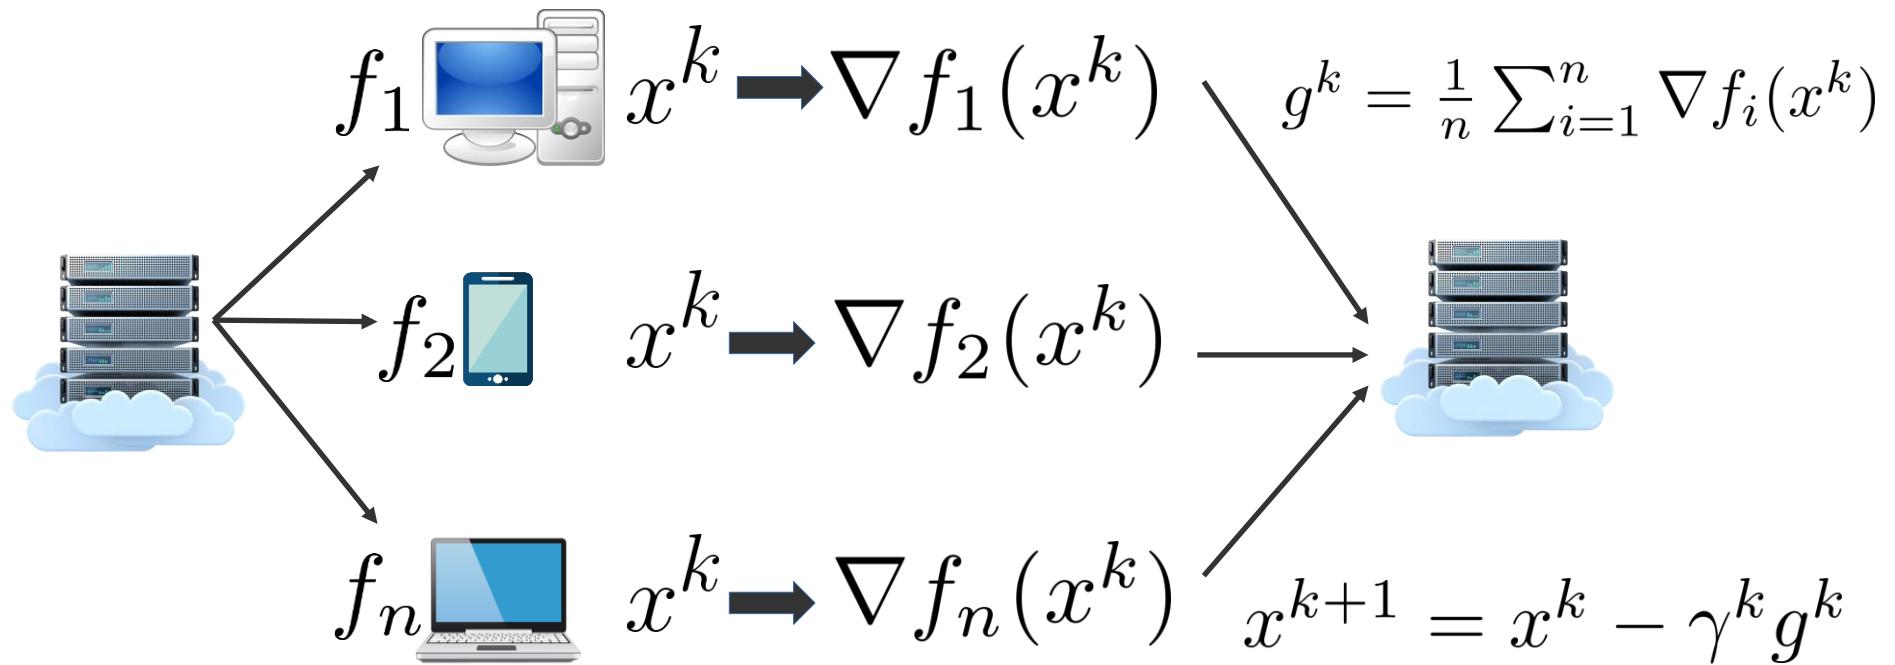
Ch8: BurTorch

Burlachenko & Richtárik, 2025

Ch9: Concluding Remarks: Summary and Future Research

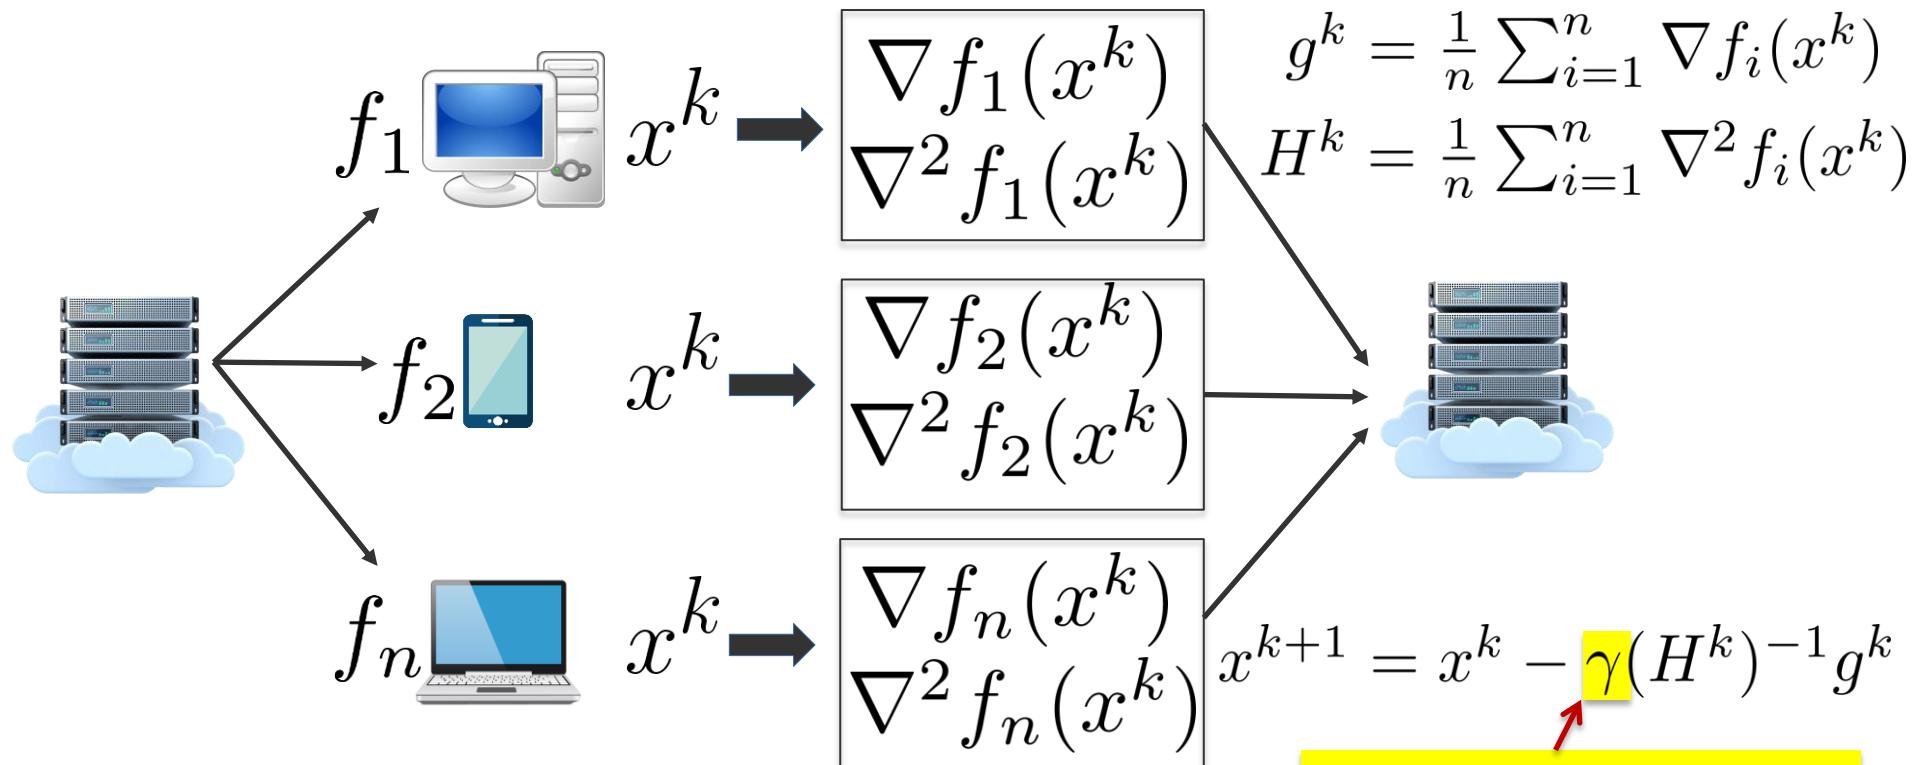
# Distributed Gradient Descent => Distributed Newton #35

• . . .



# Distributed Gradient Descent => Distributed Newton #35

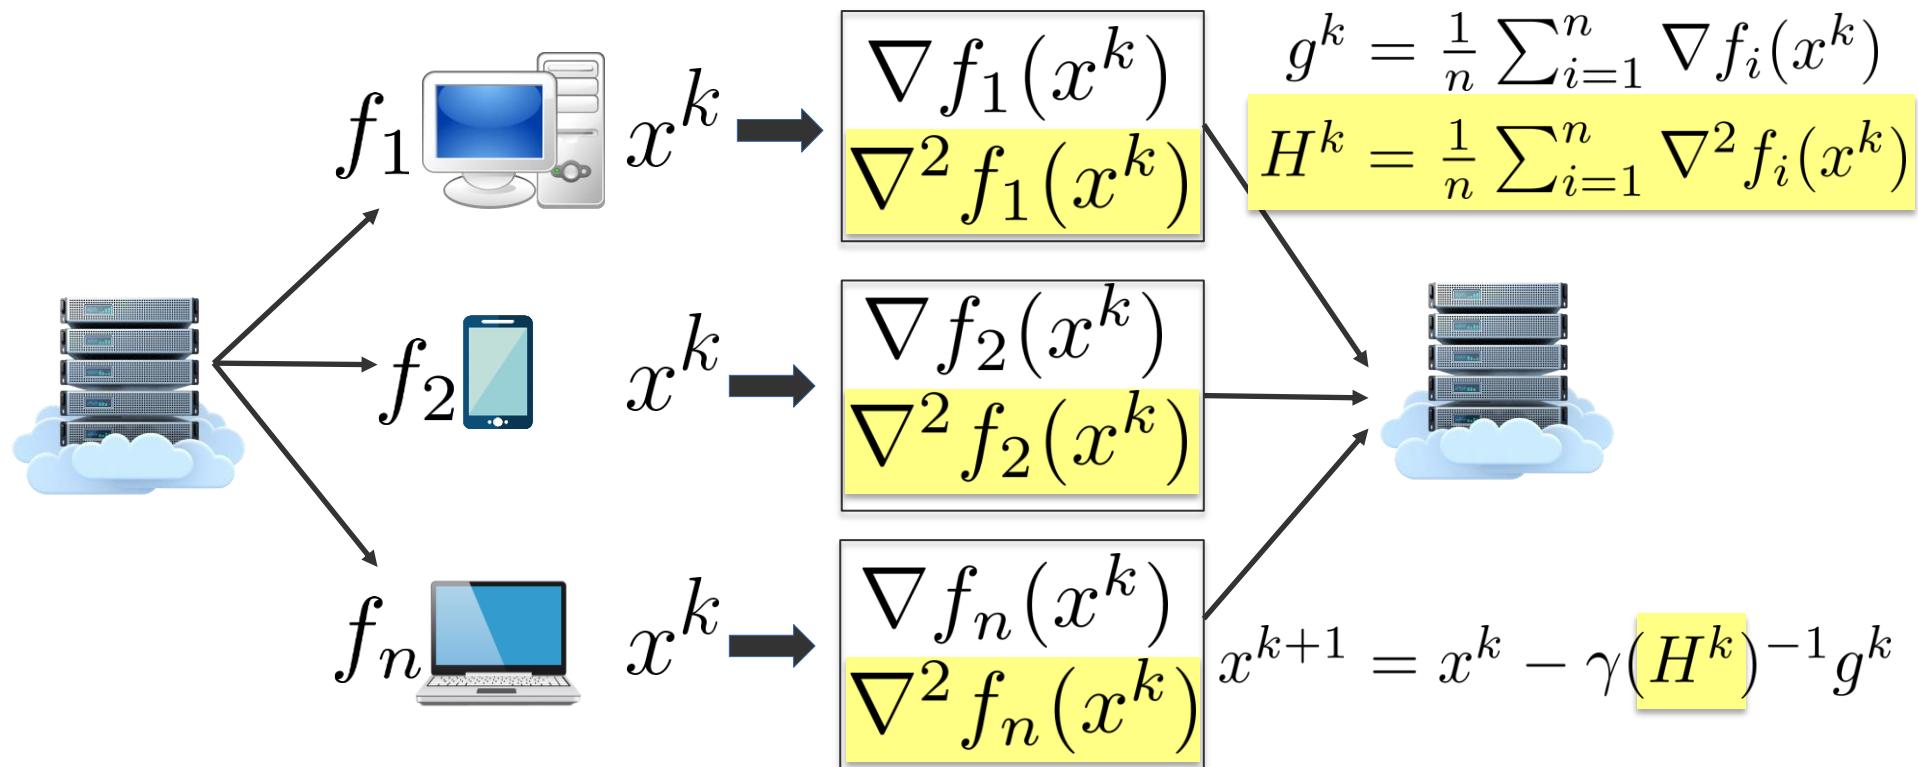
• • •



Backtracking line search  
[BV, 2004]

# Distributed Gradient Descent => Distributed Newton #35

• • •



**Bad:** The memory requirement for forming and storing second-order information

# Distributed Gradient Descent => Distributed Newton #35

$$\mu \cdot I \preceq \nabla^2 f(x) \preceq L \cdot I$$

$$\|\nabla^2 f(x) - \nabla^2 f(y)\|_2 \leq L_* \|x - y\|_2$$

$$\|\nabla f(x^k)\| \leq \eta \implies \gamma = 1, \quad \frac{L_*}{2\mu^2} \|\nabla f(x^{k+1})\|_2 \leq \left( \frac{L_*}{2\mu^2} \|\nabla f(x^k)\|_2 \right)^2$$

[BV, 2004]



$x^k \rightarrow$

$$\begin{bmatrix} \nabla f_n(x^k) \\ \nabla^2 f_n(x^k) \end{bmatrix}$$

$$x^{k+1} =$$

$$x^k - \gamma(H^k)^{-1}g^k$$

**Good:**  $\text{Error}^{k+1} \leq \text{Const} \cdot (\text{Error}^k)^2$ . In practice, number of iterations is  $\approx 6$ .

# Federated Newton Learn (2022) (Existing)

#36  
• • •

## Problem

$$x^* = \underset{x \in \mathbb{R}^d}{\operatorname{argmin}} \left( f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x) \right)$$

## FedNL: Federated Newton Learn (M. Safaryan et al., 2022)

$$\begin{aligned} x^{k+1} &= x^k - \left[ \frac{1}{n} \sum_{i=1}^n H_i^k + l^k I \right]^{-1} \frac{1}{n} \sum_{i=1}^n \nabla f_i(x^k) \\ H_i^{k+1} &= H_i^k + C_i^t (\nabla^2 f_i(x^{k+1}) - H_i^k) \end{aligned}$$

EF21 Mechanism

FedNL Technicality

Generalizing Biased and Unbiased Compressors to Symmetric Matrices

## Assumptions for FedNL Family

$f(x)$  is  $\mu$  strongly convex and  $f_i(x)$  has Lipschitz continuous Hessian

# Federated Newton Learn (2022) (Existing)

#36  
• • •

## Problem

$$x^* = \underset{x \in \mathbb{R}^d}{\operatorname{argmin}} \left( f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x) \right)$$

## FedNL: Federated Newton Learn (M. Safaryan et al., 2022)

$$\begin{aligned} x^{k+1} &= x^k - \left[ \frac{1}{n} \sum_{i=1}^n H_i^k + l^k I \right]^{-1} \frac{1}{n} \sum_{i=1}^n \nabla f_i(x^k) \\ H_i^{k+1} &= H_i^k + C_i^t (\nabla^2 f_i(x^{k+1}) - H_i^k) \end{aligned}$$

## FedNL Local Superlinear Convergence Guarantees

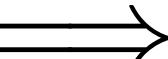
$D$  depend on Smoothness of Hessian

$C \in \{2, 8\}$

For  $\|C(M) - M\|_F^2 \leq (1 - \alpha)\|M\|_F$  the  $A = \frac{\alpha}{4}$

$$\|x^0 - x^*\| \leq \frac{\mu}{\sqrt{2D}}$$

$$H^k \leq \frac{\mu^2}{4C}$$



$$\|x^k - x^*\|^2 \leq \frac{1}{2^k} \|x^0 - x^*\|^2$$

$$\mathbf{E} \left[ \frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \right] \leq \left( 1 - \min \left( \frac{1}{3}, A \right) \right)^k \frac{\Phi^0}{\mu^2} \cdot c$$

$$H^k := \frac{1}{n} \sum_{i=1}^n \|H_i^k - \nabla^2 f_i(x^*)\|_F^2$$

$$\|\nabla^2 f(x) - \nabla^2 f(y)\|_F \leq L_F \|x - y\|_2$$

$$\Phi^k := H^k + 6BL_F^2 \|x^k - x^*\|^2$$

$$\mathbf{E}[\Phi^k] \leq \left( 1 - \min \left( \frac{1}{3}, A \right) \right)^k \Phi^0$$

# Federated Newton Learn (2022) (Existing)

#36

• • •

Can we practically use the FedNL implementation presented at ICML 2022?

# Federated Newton Learn (2022) (Existing)

#36  
• • • •

Can we practically use the FedNL implementation presented at ICML 2022?

**Not Yet!**

Requires 4.8 hours to launch a single experiment on a server-grade workstation

The prototype supports only a multi-node simulation

Prototype integration into resource-constrained applications is challenging

## Fact #1 from Data-intensive Computing and OS

Overheads introduced from a **large** number of components lead to degradation  
[1] Scalability! but at what COST?" by McSherry et.al., 2015

## Fact #2 from Computer Architecture/Programming Languages

Scripting languages offer the advantage of democratizing implementations.  
But their **eco-system clashes too much with the principles of the real hardware.**  
[2] *There's plenty of room at the top: What will drive computer performance after Moore's law?* by Leiserson et al., Science 2020

Fact #3 from S. Boyd: *This is awesome! I want to use whatever you so-called Control. What should I use? ...nothing.*

[3] InControl podcast: Interview with Stephen Boyd, 2023, 01:17:10

Fact #4 from P. Liang: *We are in a crisis. Researchers are disconnected from underlying technologies through abstractions. The problem is abstractions are leaky.*

[4] P. Liang, Stanford CS336 Language Modeling from Scratch, Spring 2025

Simultaneously advancing a **rigorous theoretical framework** and an **efficient implementation** presents a significant challenge, as both are equally demanding

## Contributions for Making FedNL Practical

1. We proposed two new practical compressors
2. Reduced wall clock time of a baseline by **×1000**
3. Outperforms several best practice solutions
4. Complete independence on 3rd party frameworks  
[Linux, Win, macOS] x [AArch64, x86-64, CUDA]
5. Can be utilized as native OS executable binaries and libraries

# Pessimism of TopK Contraction Factor

#39

$$\mathcal{C}^d(\alpha) = \{C \mid C : \mathbb{R}^d \rightarrow \mathbb{R}^d, \quad \mathbb{E} [\|C(x) - x\|^2] \leq (1 - \alpha)\|x\|^2\}, \quad \forall \alpha \in (0, 1]$$

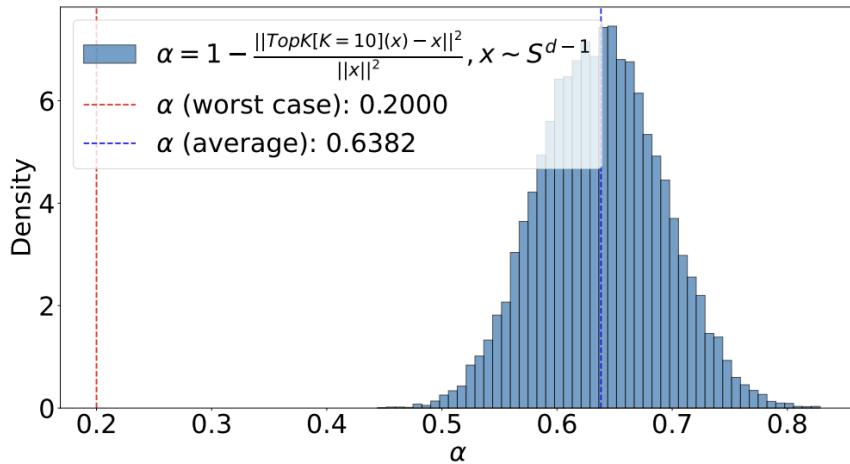
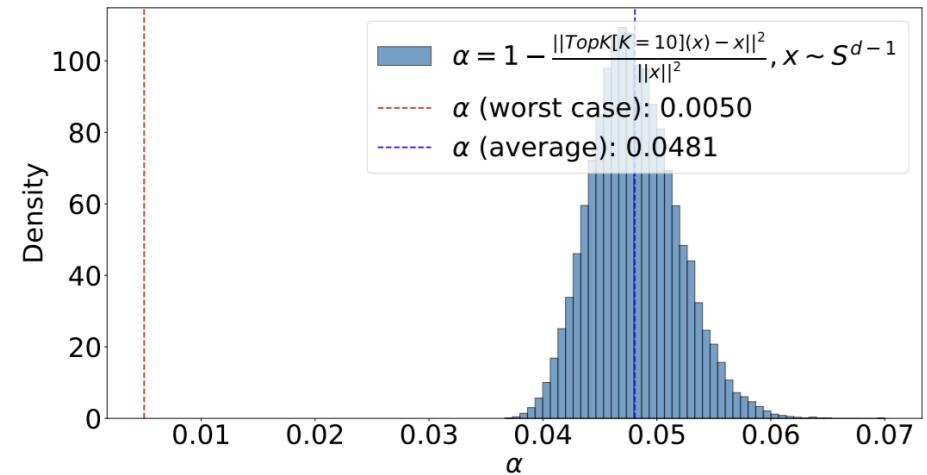
(a)  $d=50$ , **TopK** with  $k = 10$ .(b)  $d=2000$ , **TopK** with  $k = 10$ .

Figure 7.1: Discrepancy between worst-case  $\alpha$  and  $\alpha(x)$  when  $x \sim_{\text{u.a.r.}} S^{d-1}$ . Number of trials 20 000.

# Adaptive TopLEK Compressor

$$\mathbb{E} [\|TopLEK(x) - x\|^2] = (1 - \alpha)\|x\|^2, \quad 0 < \alpha \leq 1$$

- The idea is to perform compression using **TopK**, with smaller parameter  $\hat{k} \leq K$  compressing as much as theoretically allowed — but no more
- For **TopLEK**, the inequality that describes contraction becomes a tight equality

$$\mathcal{B}^d(\omega) = \{B \mid B : \mathbb{R}^d \rightarrow \mathbb{R}^d, \quad \mathbb{E} [\|B(x) - x\|^2] \leq \omega \|x\|^2, \quad \mathbb{E} [B(x)] = x\}, \quad \forall \omega \geq 0$$

**RandK** selects a subset of coordinates of cardinality  $k$  u.a.r. from a total of  $d$  coordinates, zeroing out the rest and scaling the output to preserve unbiasedness

# Cache-aware RanSeqK Compressor

#41  
• •

$$\mathcal{B}^d(\omega) = \{B \mid B : \mathbb{R}^d \rightarrow \mathbb{R}^d, \quad \mathbb{E} [\|B(x) - x\|^2] \leq \omega \|x\|^2, \quad \mathbb{E} [B(x)] = x\}, \quad \forall \omega \geq 0$$

**RanSeqK** selects a pivot  $s$  u.a.r. from  $\{1, 2, \dots, d\}$ , then selects deterministically a block of size  $k$  from  $(1, 2, \dots, d)$  seen as a torus

It has the same variance as **RandK**,  
but it is more appealing in practice

Table 7.8: Memory latency comparison in computing devices.

Device and Memory Level	Approximate Latency (ns)	Scale
CPU cycle	0.3	x1
CPU register (SRAM)	0.3	x1
L1 cache (SRAM)	0.9	x3
Floating Point addition, subtraction, and multiplication	1.2	x4
L2 cache (SRAM)	3	x10
L3 cache (SRAM)	10	x33
Main memory or Physical Memory (DRAM)	100	x330
The OS System Call: Transitioning from user to kernel space	300	x1000
Solid-State Disk (SSD)	10 000	x33 000
Rotational Hard Disk Drive (HDD)	10 000 000	x33 000 000

# Ch7: Unlocking FedNL

#42

## Single Node Experiments: L2 Regularized Logistic Regression

### Baseline Improvements

Table 7.1: Single-node setting,  $n = 142$ , FedNL (B),  $r = 1000$ ,  $\lambda = 0.001$ ,  $\alpha$  - option 2, FP64, 24 cores at 3.3 GHz.

#	Client Compression	$\ \nabla f(x^{last})\ $	Total Time (seconds)
1	RandK[K=8d] (We)	$3 \cdot 10^{-18}$	18.84
2	RandK[ $k = 8d$ ] (Base)	$3 \cdot 10^{-18}$	17510.00
3	TopK[K=8d] (We)	$2.80 \cdot 10^{-18}$	18.72
4	TopK[ $k = 8d$ ] (Base)	$2.80 \cdot 10^{-18}$	19770.00
5	RandSeqK[K=8d] (We)	$3.19 \cdot 10^{-18}$	16.70
6	TopLEK[K=8d] (We)	$3.45 \cdot 10^{-18}$	18.48
7	Natural (We)	$3.10 \cdot 10^{-18}$	27.02
8	Ident (We)	$2.46 \cdot 10^{-18}$	24.12

### Single Node

Table 7.2: Single-node setting,  $n = 142$ , FedNL-LS (B),  $\|\nabla f(x^{last})\| \approx 9 \cdot 10^{-10}$ , FP64, 24 cores at 3.3 GHz.

#	Solver	W8A, $d = 301$ , $n_i = 350$	A9A, $d = 124$ , $n_i = 229$	PHISHING, $d = 69$ , $n_i = 77$
Initialization Time (seconds)				
1	CVXPY	+2.54	+2.33	+2.28
2	FedNL	+0.939	+0.196	+0.081
Solving Time (seconds)				
3	CLARABEL	19.24	10.83	2.50
4	ECOS	22.22	8.02	2.55
5	ECOS-BB	22.00	8.00	2.12
6	SCS	31.14	19.36	4.57
7	MOSEK	16.90	9.59	3.55
8	FedNL-LS / RandK[ $k = 8d$ ]	4.35	0.34	0.12
9	FedNL-LS / RandSeqK[ $k = 8d$ ]	3.34s	0.29	0.06
10	FedNL-LS / TopK[ $k = 8d$ ]	4.49	0.46	0.10
11	FedNL-LS / TopLEK[ $k = 8d$ ]	4.79	0.34	0.61
12	FedNL-LS / Natural	3.13	0.17	0.08
13	FedNL-LS / Identical	0.63	0.09	0.06

# Ch7: Unlocking FedNL

## Multi Node Experiments: L2 Regularized Logistic Regression

Table 7.3: Multi-node setting,  $n = 50$  clients, 1 master,  
 $|\nabla f(x^{last})| \approx 10^{-9}$ , FP64, 1 CPU core/node.

#	Solution	W8A $d = 301,$ $n_i = 994$	A9A $d = 124,$ $n_i = 651$	PHISHING $d = 69,$ $n_i = 221$
<b>Initialization Time (seconds)</b>				
1	Ray		+52.0	
2	Apache Spark		+25.82	
3	<b>FedNL</b>		+1.1	
<b>Solving Time (seconds)</b>				
4	Ray	116.17	28.13	11.54
5	Apache Spark	36.65	33.59	33.14
6	<b>FedNL</b> / RandK[ $k = 8d$ ]	12.6	4.52	0.21
7	<b>FedNL</b> / RandSeqK[ $k = 8d$ ]	12.56	5.10	0.14
8	<b>FedNL</b> / TopK[ $k = 8d$ ]	12.20	5.79	5.23
9	<b>FedNL</b> / TopLEK[ $k = 8d$ ]	15.11	3.26	0.82
10	<b>FedNL</b> / Natural	5.75	1.56	0.14

# Ch7: Structure of x1000 Time Improvement for Single Node

#44  
FP64, 3.3Ghz, 12 cores, Intel(R) CPU. Logistic Regression d=301

Baseline: Single node Python/NumPy implementation from the original paper	x1
Rewrite in pure C++20/CMake with support macOS, Linux, Windows	x20
Data Processing Optimization	x1.077
Eliminating Integer Division	x1.225
Utilizing AVX512 CPU Extension in x86-64	x1.379
Compiler and Linker Optimization	x1.128
Use Sparsity from FedNL Compressors	x1.44
Reuse Computation from Oracles	x1.50
Basic Linear Algebra Improvements	x1.338
Linear System Solve Improvement	x1.31
Custom oracles without using Cache-Oblivious matrix multiplication	x3.072
Better Compressors Implementation:	x1.14
Thread pool of workers equal to number of physical cores, atomics for sync	x1.412
Mem. Optimization. Custom client-specific memory pool instead of global	x1.278

# Federated Learning Challenges Addressed in the Thesis #45

## Theoretical Work

1. Data Heterogeneity

2. Device Heterogeneity

3. Communication Bottleneck

4. Privacy

5. Software

## Theory-Inspired Practical Work

## Practical Work

### Ch1: Introduction



Ch3: EF21-W  
Richtárik et al., 2024

Ch3: EF21-W  
Richtárik et al., 2024

Ch4: DCGD/PERMK/AES  
Burlachenko et al., 2023

Ch2: FL\_PyTorch  
Burlachenko et al., 2021

Ch5: PAGE Extensions  
Tyurin et al., 2023



Unlocking FedNL  
Richtárik, 2024

Ch7: Unlocking FedNL  
Burlachenko and Richtárik, 2024

Ch8: BurTorch

Burlachenko & Richtárik, 2025

Ch8: BurTorch

Burlachenko & Richtárik, 2025

Ch9: Concluding Remarks: Summary and Future Research

# Federated Learning Challenges Addressed in the Thesis #45

## Theoretical Work

1. Data Heterogeneity

2. Device Heterogeneity

3. Communication Bottleneck

4. Privacy

5. Software

## Theory-Inspired Practical Work

## Practical Work

### Ch1: Introduction

Ch3: EF21-W  
Richtárik et al., 2024

Ch3: EF21-W  
Richtárik et al., 2024

Ch4: DCGD/PERMK/AES  
Burlachenko et al., 2023

Ch5: PAGE Extensions  
Tyurin et al., 2023

Ch6:  
Compressed L2GD  
Bergou et al., 2023

Ch2: FL\_PyTorch  
Burlachenko et al., 2021

Ch7: Unlocking FedNL  
Burlachenko and Richtárik, 2024

Ch8: BurTorch  
Burlachenko & Richtárik, 2025

Ch7: Unlocking FedNL  
Burlachenko & Richtárik, 2024

Ch8: BurTorch  
Burlachenko & Richtárik, 2025

Ch9: Concluding Remarks: Summary and Future Research

# Chapters Excluded from In-Depth Presentation

#46

**Ch2: FL\_Pytorch (2021)** Existing software frameworks for FL prioritize deployment, raise the entry barrier, and demand expertise in distributed systems. Research requires tools with functionalities distinct from industrial runtimes.



**Ch5: PAGE Extensions (2023)** PAGE is a theoretical optimal algorithm for finding a stationary point in sampled gradient complexity in big -  $\mathcal{O}$  notation. This work enhances the analysis of PAGE and extends it with other sampling strategies.



**Ch6: Compressed L2GD (2023)** New Paradigm for FL was proposed in “*Federated Learning of a Mixture of Global and Local Models*” (2021) by F.Hanzely and P.Richtárik. This work extends it with bidirectional communication compressors.



**Ch8: BurTorch (2025)** Latency-efficient backpropagation CPU implementation, which outperforms: JAX, TF, TF Lite, LibTorch (C++), PyTorch TorchScript, PyTorch Python, Apple MLX, Autograd, Micrograd in memory, time, consumed energy.

UNDER REVIEW



# Thank You for Your Time and Attention!

All presented projects are accompanied by open-source code,  
promoting a culture of openness and collaboration

# Backup Slides



جامعة الملك عبد الله  
للعلوم والتكنولوجيا  
King Abdullah University of  
Science and Technology

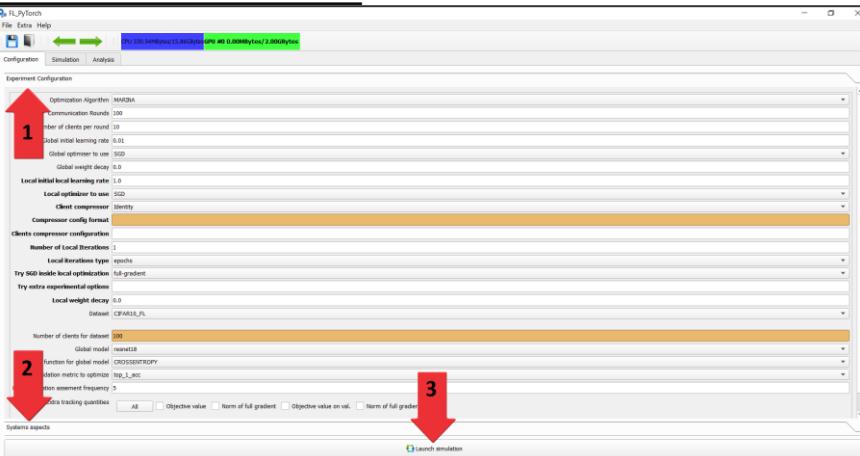
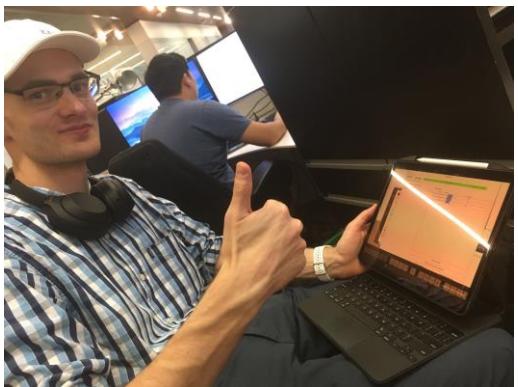


TOOL	Docs.	DX	Lang.	AI frameworks	AI type	Examples	Dist. channels	Multinode?	C/V	H/V	Sync/Aync	PVC/SEC	Tools integration	TOTAL
Flower [49]	1.5	1.5	2	2.5	2	1	1	1	1	0	1	1	1.5	17
OpenFL [50]	1.5	1.5	2	2	1	1	1	1	1	1	1	1	1.5	16.5
IBM-Federated [45]	1.5	1	1	2	2	1	1	1	1	0	0	1	1.5	14
PySyft [46]	1.5	1	1	2	1	1	1	1	1.5	0	0	1	1.5	13.5
Nvidia Flair [43]	1	1	1.5	2	2	1	1	1	1	0	0	1	1	13.5
FedML [48]	1.5	0	1	2.5	2	1	1	1	1	0	1	0	1.5	13.5
Fedn [60]	1	1	2	2	1	1	0	1	0.5	1	0	1	1.5	13
FedLearn-algo [54]	0	0	1.5	2	2	1	0	1	1	1	0	1	1.5	12
XFL [81]	1	0	1	2	2	1	0	1	1	1	0	1	1	12
PLATO [80]	1	0	1	2.5	1	1	1	1	1	0	1	0	1.5	12
FATE [47]	1	0	1.5	0	2	1	0	1	1	1	1	1	1.5	12
APPFL [62]	1	0	2	1	1	1	1	1	1	0	1	0	1	11
FedLab [51]	1	0	2	1	1	1	1	1	1	0	1	1	0	11
FedBioMed [61] (GitLab)	0	1	1	2	1	1	0	1	0.5	0	0	1	1.5	10
FedJAX [55]	1	1.5	1	2	1	1	1	1	0.5	0	0	1	0	10
OpenFED [37]	1	1	2	1	1	1	1	1	0.5	0	0	1	0	9.5
Tensorflow Federated [44]	1	1	2	1	1	1	1	1	0.5	0	0	1	0	9.5
PyVertical [56] [57]	0	1	2	0	1	1	0	1	1	0	0	1	1	9
FL-Pytorch [71]	1.5	1	1	1	1	1	1	1	0.5	0	0	0	0	8
FLUTE [79]	0	0	1	1	1	1	0	0.5	0	0	1	0	1	7.5
PriMIA [58]	1	0	0	1	1	1	0	0.5	0	0	0	1	1.5	7
Sunday FL [66]	1	0	1.5	0	0	1	0	1	1	0	0	0	1	6.5
dsMTL [65]	0.0	0.0	1.0	1.0	1.0	1.0	0.0	1.0	0.0	0.0	0	1.0	0.0	6
Substra [59]	1	0	0	0	0	1	1	1	1	0	1	0	0	6
DecFL [67]	0	0	0	1	1	1	0	1	0.5	0	0	1	0	5.5
Vantage6 [69]	1.5	0	1	0	0	0	1	1	1	0	0	0	0	5.5
HyFed [63]	0	1	0	0	0	1	0	1	0	0	0	1	0	4
MTC-ETH [68]	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.5	0	0	1.0	0.0	2.5

## FeLebrities (2023):

### A User-Centric Assessment of Federated Learning Frameworks

W. Riviera,  
I.B. Galazzo,  
G. Menegaz



# Ch5: PAGE Extensions

#49

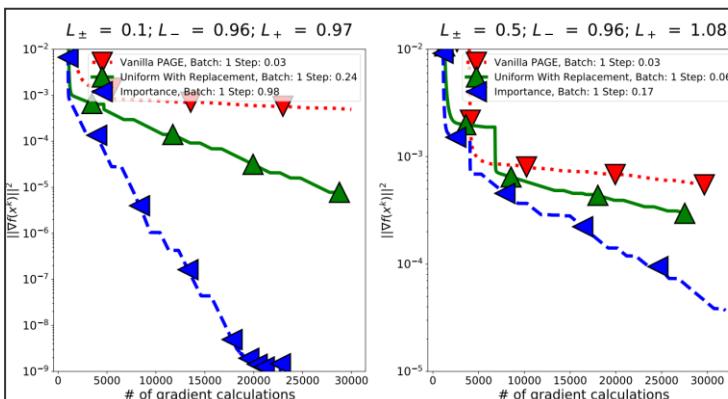
## Key Step in Discovery Path

The  $\Omega$  is the probability sample space. Let  $S : a_1 \times \dots \times a_n \times \Omega \rightarrow m$  with the following properties:

1.  $\mathbb{E}_w[S(\cdot)] = \frac{1}{n} \sum_{i=1}^n a_i$
2.  $\mathbb{E}_w [\|S(\cdot) - \frac{1}{n} \sum_{i=1}^n a_i\|^2] = \frac{A}{n} \sum_{i=1}^n \left( \frac{1}{nw_i} \|a_i\|^2 \right) - B \|\frac{1}{n} \sum_{i=1}^n a_i\|^2$

In the paper, we provide several sampling strategies that satisfy the above conditions. Next, with respect to the weighting  $w_i$ , we require to estimate the constants  $L_{+,w}$  and  $L_{\pm,w}$ :

1.  $\frac{1}{n} \sum_{i=1}^n \frac{1}{nw_i} \|\nabla f_i(x) - \nabla f_i(y)\|^2 \leq L_{+,w}^2 \|x - y\|^2$
2.  $\frac{1}{n} \sum_{i=1}^n \frac{1}{nw_i} (\|\nabla f_i(x) - \nabla f_i(y)\|^2 - \|\nabla f(x) - \nabla f(y)\|^2) \leq L_{\pm,w}^2 \|x - y\|^2$



## Theoretical Improvements

Our Theory for Single Gradient Oracles:

$$N = \mathcal{O} \left( n + \frac{\Delta_0}{\epsilon^2} \cdot |S| \left( L_- + \sqrt{\frac{n}{|S|} ((A - B)L_{+,w}^2 + BL_{\pm,w}^2)} \right) \right)$$

Original PAGE Analysis:

$$N_{\text{orig}} = \mathcal{O} \left( n + \frac{\Delta_0}{\epsilon^2} \cdot (L_- + \sqrt{n}L_+) \right)$$

Improved Original PAGE Analysis:

$$N_{\text{improved}} = \mathcal{O} \left( n + \frac{\Delta_0}{\epsilon^2} \cdot (0 + \sqrt{n}L_+) \right)$$

Important Sampling in PAGE:

$$N_{\text{important-sampling}} = \mathcal{O} \left( n + \frac{\Delta_0}{\epsilon^2} \cdot \tau \cdot \left( L_- + \sqrt{\frac{n}{\tau} L_{\pm,w}} \right) \right)$$

$$L_{\pm,w} \leq L_+ \quad \text{and} \quad L_{+,w} \leq L_+$$

Comparison on synthesized quadratics

$$\min_{x_1, \dots, x_n \in \mathbb{R}^d} \left\{ F_\lambda(x) := \left( \frac{1}{n} \sum_{i=1}^n f_i(x_i; D_i) \right) + \lambda \cdot \frac{1}{2n} \sum_{i=1}^n \|x_i - \bar{x}\|^2 \right\}$$

$$f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x_i; D_i) \quad h(x) := \lambda \cdot \frac{1}{2n} \sum_{i=1}^n \|x_i - \bar{x}\|^2$$

**Goals**

- Strongly convex case  $\mathbb{E} [x^k - x^*] \leq \varepsilon \|x^0 - x^*\|^2$
- Non-convex case  $(\mathbb{E} [\|\nabla F(x^k)\|])^2 \leq \mathbb{E} [\|\nabla F(x^k)\|^2] \leq \varepsilon^2$

**Compressors**

Each client and master use its own unbiased compressor from  $\mathcal{B}^d(w)$

**Aspects of L2GD**

- $\lambda$  is scalar parameter which allows tradeoff between local and global model  
 $\lambda \rightarrow 0$  ( $\lambda \rightarrow +\infty$ ) all nodes are working in decoupled (coupled) form
- $p$  is probability of making (relaxed) aggregation
- $1 - p$  is a probability to make in parallel local GD steps

**Results**

- The work extends paper [HR, 2021] with bidirectional unbiased compressors
- Linear convergence rate to neighborhood (strongly convex case) and theory for non-convex case
- Optimal values of  $p(\lambda, L := nL_f)$
- Extended empirical study
- Highlighted that **FedAVG** is a particular case of **L2GD** when  $\eta\lambda \approx np$

**Algorithm 15** Compressed L2GD.

---

```

Input: step size  $\eta > 0$ , probability  $p$ 
Initialize:  $\{x_i^0\}_{i=1,\dots,n}$ ,  $\xi_{-1} = 1$ ,  $\bar{x}^{-1} = \frac{1}{n} \sum_{i=1}^n x_i^0$ 
for  $k = 0, 1, 2, \dots$  do
  Draw:  $\xi_k = 1$  with probability  $p$ 
  if  $\xi_k = 0$  then
    on all devices:  $x_i^{k+1} = x_i^k - \frac{\eta}{n(1-p)} \nabla f_i(x_i^k)$  for  $i \in [n]$ 
  else
    if  $\xi_{k-1} = 0$  then
      on all devices: Compress  $x_i^k$  to  $\mathcal{C}_i(x_i^k)$  and communicate  $\mathcal{C}_i(x_i^k)$  to the master
      on master:
        1. Receive  $\mathcal{C}_i(x_i^k)$  from all devices  $i \in [n]$ 
        2. Compute  $\bar{y}^k \stackrel{\text{def}}{=} \frac{1}{n} \sum_{j=1}^n \mathcal{C}_j(x_j^k)$ 
        3. Compress  $\bar{y}^k$  to  $\mathcal{C}_M(\bar{y}^k)$ 
        4. Communicate  $\mathcal{C}_M(\bar{y}^k)$  to all devices
    on all devices: Perform aggregation step  $x_i^{k+1} = x_i^k - \frac{\eta\lambda}{np} (x_i^k - \mathcal{C}_M(\bar{y}^k))$ 
  else
    on all devices:
      a.  $\bar{x}^k = \bar{x}^{k-1}$ 
      b. Perform aggregation step  $x_i^{k+1} = x_i^k - \frac{\eta\lambda}{np} (x_i^k - \bar{x}^k)$ 
    end if
  end if
end for

```

Table 6.2: Summary of the benchmarks. The measured quantity is bits/n to achieve 0.7 Top1 test accuracy, with  $n = 10$  clients. For DenseNet-121, MobileNet, ResNet-18 the baseline is FedAVG with Natural compressor with 1 local epoch.

Model	Training Parameters	L2GD bits/n	Baseline bits/n
DenseNet-121	$79 \times 10^5$	$8 \times 10^{11}$	$4 \cdot 10^{15}$
MobileNet	$32 \times 10^5$	$1.7 \times 10^{11}$	$1 \times 10^{15}$
ResNet-18	$11 \times 10^6$	$1.1 \times 10^{12}$	$1.5 \times 10^{16}$



# Ch8: BurTorch

#51

## Benchmarks Across Linux, macOS, and Windows

Table 8.2: Backpropagation over 100K iterations with a tiny compute graph from Figure 8.1. Mean and std. deviation across 5 launches, FP64, Windows OS. See also Figure 8.3. The numerical results across frameworks match exactly.

#	Framework, Mode, Language	Device	Compute Time (sec.)	Relative to BurTorch
1	BurTorch, Eager, C++	CPU	$0.007 \pm 0.0004$	$\times 1.0$ (We)
2	TensorFlow 2.8.0, Eager, Python	CPU	$55.217 \pm 0.2975$	$\times 7888.1$
3	TensorFlow 2.8.0, Graph, Semi-Python	CPU	$14.469 \pm 0.0734$	$\times 2067.0$
4	TF Lite 2.8.0, Graph, TF Lite Interpreter	CPU	$0.589 \pm 0.0102$	$\times 84$
5	Autograd 1.7.0, Eager, Python	CPU	$18.956 \pm 0.2962$	$\times 2708.0$
6	PyTorch 2.5.1, Eager, Python	GPU	$51.380 \pm 0.4666$	$\times 7340.0$
7	PyTorch 2.5.1, Eager, Python	CPU	$10.419 \pm 0.0647$	$\times 1488.4$
8	PyTorch 2.5.1, Graph, TorchScript	CPU	$9.994 \pm 0.1021$	$\times 1428.5$
9	PyTorch 2.5.1, Eager, LibTorch, C++	CPU	$5.300 \pm 0.0667$	$\times 757.14$
10	JAX 0.4.30, Eager, Python	CPU	$291.764 \pm 8.5373$	$\times 41860.5$
11	JAX 0.4.30, Graph, Semi-Python	CPU	$5.580 \pm 0.0661$	$\times 797.1$
12	Micrograd, Eager, Python	CPU	$1.590 \pm 0.0152$	$\times 227.1$
13	In Theory for this CPU (Registers Only)	CPU	$\Omega(0.0004)$	$\times 0.057$

Table 8.7: BurTorch and PyTorch in training GPT-3 like model, FP32, 1 CPU core, Peak private virtual memory. Trainable variables: 46K.

Batch	BurTorch, Eager, C++	PyTorch, Graph, TorchScript		PyTorch, Eager, Python	
		Compute (ms)	Mem. (MB)	Compute (ms)	Mem. (MB)
1	$0.515 \pm 0.067$	16.7	$11.119 \pm 48.118$	1624	$11.715 \pm 10.741$
2	$1.027 \pm 0.091$	16.7	$11.177 \pm 37.138$	1623	$12.166 \pm 11.461$
4	$2.106 \pm 0.130$	16.7	$11.762 \pm 37.171$	1624	$12.424 \pm 11.120$
8	$4.222 \pm 0.238$	16.7	$12.041 \pm 36.312$	1631	$13.167 \pm 11.613$
16	$8.358 \pm 0.644$	16.7	$13.451 \pm 37.415$	1633	$14.111 \pm 11.278$
32	$16.787 \pm 1.03601$	16.7	$16.048 \pm 36.460$	1632	$16.661 \pm 11.122$
64	$31.696 \pm 0.737$	16.8	$21.794 \pm 37.302$	1640	$22.189 \pm 11.531$

Table 8.5: Comparison of BurTorch and PyTorch performance for training MLP-like model. Batch:  $b = 1$ , Compute: FP32, Single CPU core. Initialization time is end-to-end time for training with 1 iteration. Compute time excludes batch preparation. Memory is the peak private virtual memory.

#	Parameters (d) Hidden Dim.(e)	PyTorch, Eager, v2.5.1 [CPU]			BurTorch, Eager [CPU]		
		Init (ms)	Compute (ms)	Mem. (MB)	Init (ms)	Compute (ms)	Mem. (MB)
1	5,963 ( $e = 4$ )	5 540	$1.46 \pm 4.63$	2 651	15.63	$0.032 \pm 0.008$	35.8
2	18,587 ( $e = 16$ )	5 627	$1.52 \pm 4.21$	2 653	16.51	$0.074 \pm 0.016$	36.7
3	35,419 ( $e = 32$ )	5 673	$1.55 \pm 5.00$	2 653	18.24	$0.124 \pm 0.019$	38.3
4	69,083 ( $e = 64$ )	5 537	$1.63 \pm 4.62$	2 668	18.94	$0.221 \pm 0.040$	40.8
5	136,411 ( $e = 128$ )	5 799	$1.79 \pm 5.19$	2 660	21.39	$0.417 \pm 0.077$	45.9
6	540,379 ( $e = 512$ )	5 556	$3.01 \pm 5.57$	2 683	37.09	$2.093 \pm 0.429$	71.4
7	1,079,003 ( $e = 1024$ )	5 544	$5.57 \pm 6.75$	2 719	56.57	$4.550 \pm 0.847$	107.0

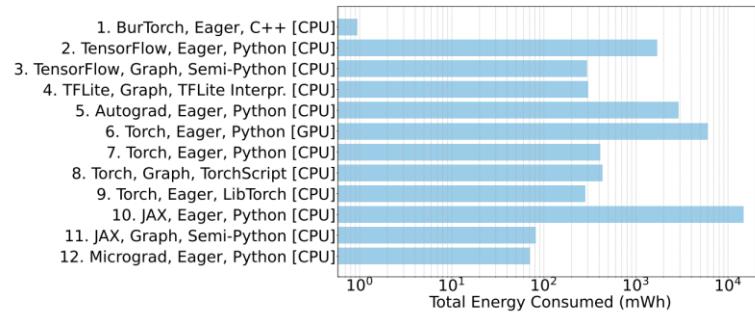
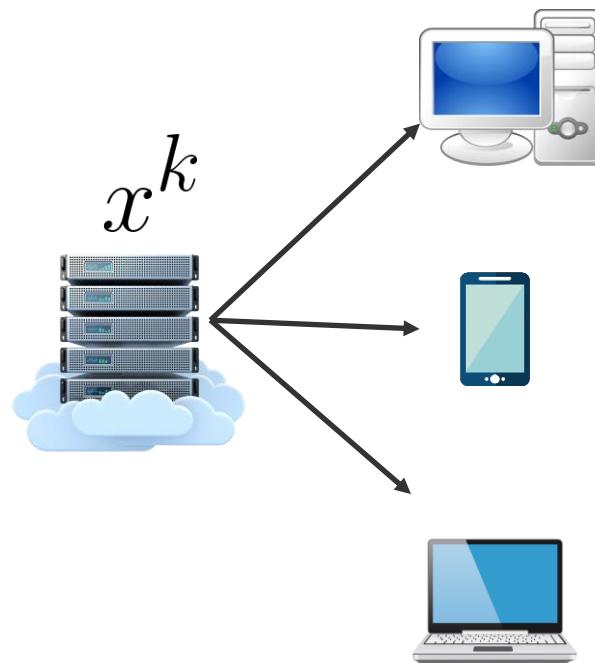


Figure 8.7: Visualization of Table 8.19. Total power drain over 200K iterations with a *small* dynamically constructed compute graph (Figure 8.2) consisting of 32 nodes, using FP64. Voltage: 11.7V, Battery: DELL J8FK941J, Chemistry: Li-poly, OS: Windows 11. The numerical results across frameworks match exactly.



Single-Node	Typical Bandwidth
CPU <-> System DDR Memory	51 200 MBytes/sec (DDR5)
GPU core <-> GPU DDR Memory	128 000 MBytes/sec (GDDR6) DRAM in NVIDIA GPU NVIDIA Ada Lovelace 1008 000 MBytes/sec
GPU DDR <-> PCI-E <-> System DDR Memory	4 000 MBytes / sec (PCI-E v5, 1 lane)
SATA 3x (HDD)	6000 Mbytes / sec
USB 3.0 (External storage)	600 MBytes / sec
GPU <-> GPU (NVLink)	50 000 Mbytes/sec
GPU <-> GPU (NVLink via NVSwitch inside DGX-2)	50 000 Mbytes/sec

Multi-Node	Typical Bandwidth
Fast Ethernet	12.5 MBytes/sec
Gigabit Ethernet	125 MBytes/sec
InfiniBand HDR	6250 Mbytes/sec
InfiniBand Mellanox	25 000 Mbytes/sec

AES is secure for encoding a **single block**  
For multiple blocks, it should be used with a  
***Mode of Operation Algorithm***

Overhead: 16 bytes/message

## AES with EAX Authentication

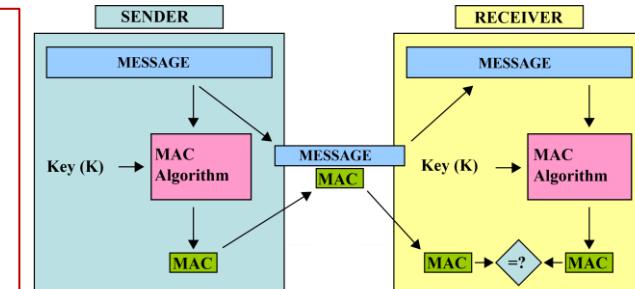
Overhead: 16 bytes/message

To ensure that message has not been altered in transit,  
AES should be paired with a  
***Message Authentication Code Algorithm***  
(Similar to CRC for non-secure applications)



Incorrect Electronic Code Book (ECB)

Images: Google Search



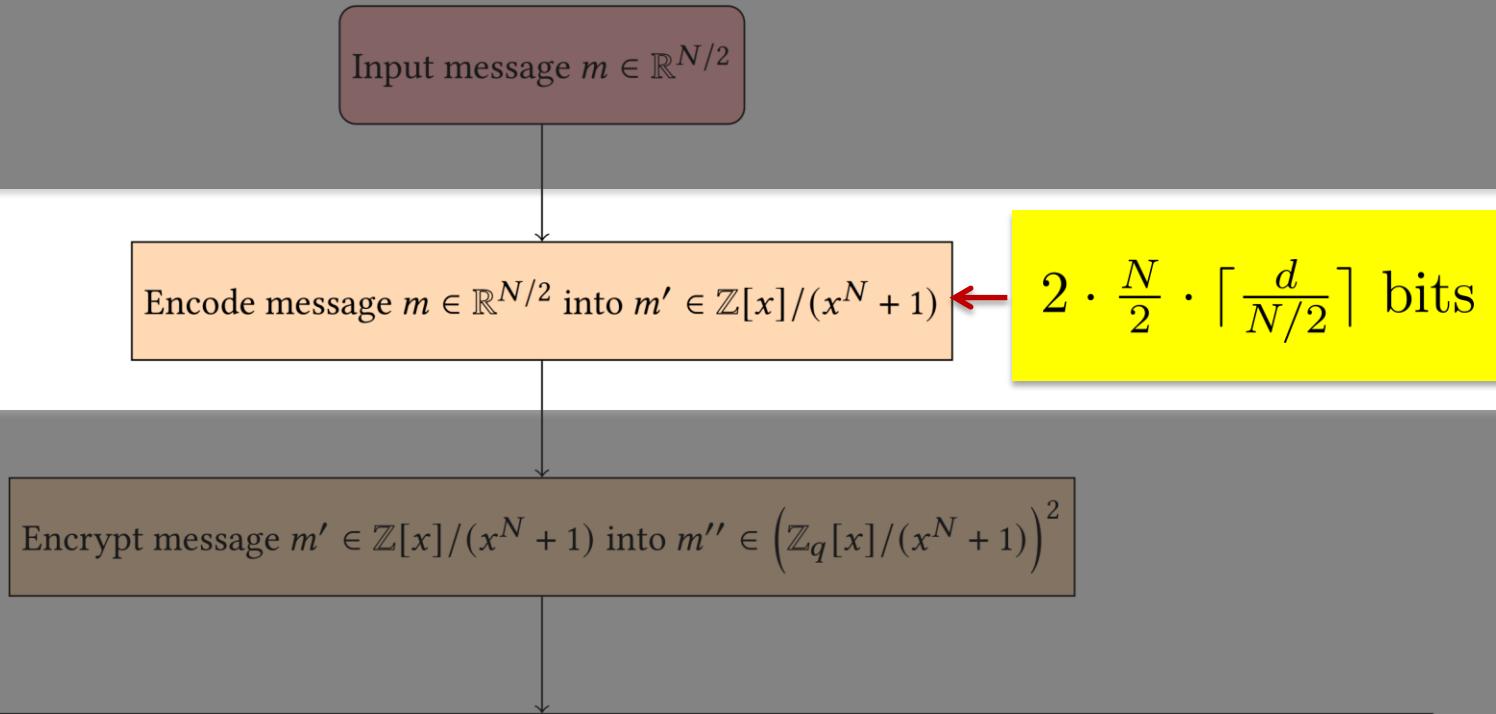
Input message  $m \in \mathbb{R}^{N/2}$

For input with  $d$  scalars, amount of ciphertexts is  $\lceil \frac{d}{N/2} \rceil$

For AES-128 compatibility level  $N > 16\,384$

Encrypt message  $m' \in \mathbb{Z}[x]/(x^N + 1)$  into  $m'' \in \left(\mathbb{Z}_q[x]/(x^N + 1)\right)^2$

Computation on encrypted messages  $(m''_1, m''_2, \dots)$ . CKKS allows to perform: addition, multiplication, and rotation.



Computation on encrypted messages  $(m''_1, m''_2, \dots)$ . CKKS allows to perform: addition, multiplication, and rotation.

# Ch4: A Day of Life for Message with CKKS Scheme #54

# bits to store  $m''$  is:  $2 \cdot 2 \cdot \frac{N}{2} \cdot \lceil \frac{d}{N/2} \rceil \cdot Q$

$$m''(\text{shared}) = pk(\text{secret}) \cdot m' + \text{noise}$$

Encode message  $m \in \mathbb{R}^{N/2}$  into  $m' \in \mathbb{Z}[x]/(x^N + 1)$

# bits to store  $q$  is  $Q = \log_2(q)$   
 $Q \geq 438$  for AES-128 level

Encrypt message  $m' \in \mathbb{Z}[x]/(x^N + 1)$  into  $m'' \in (\mathbb{Z}_q[x]/(x^N + 1))^2$

Computatio

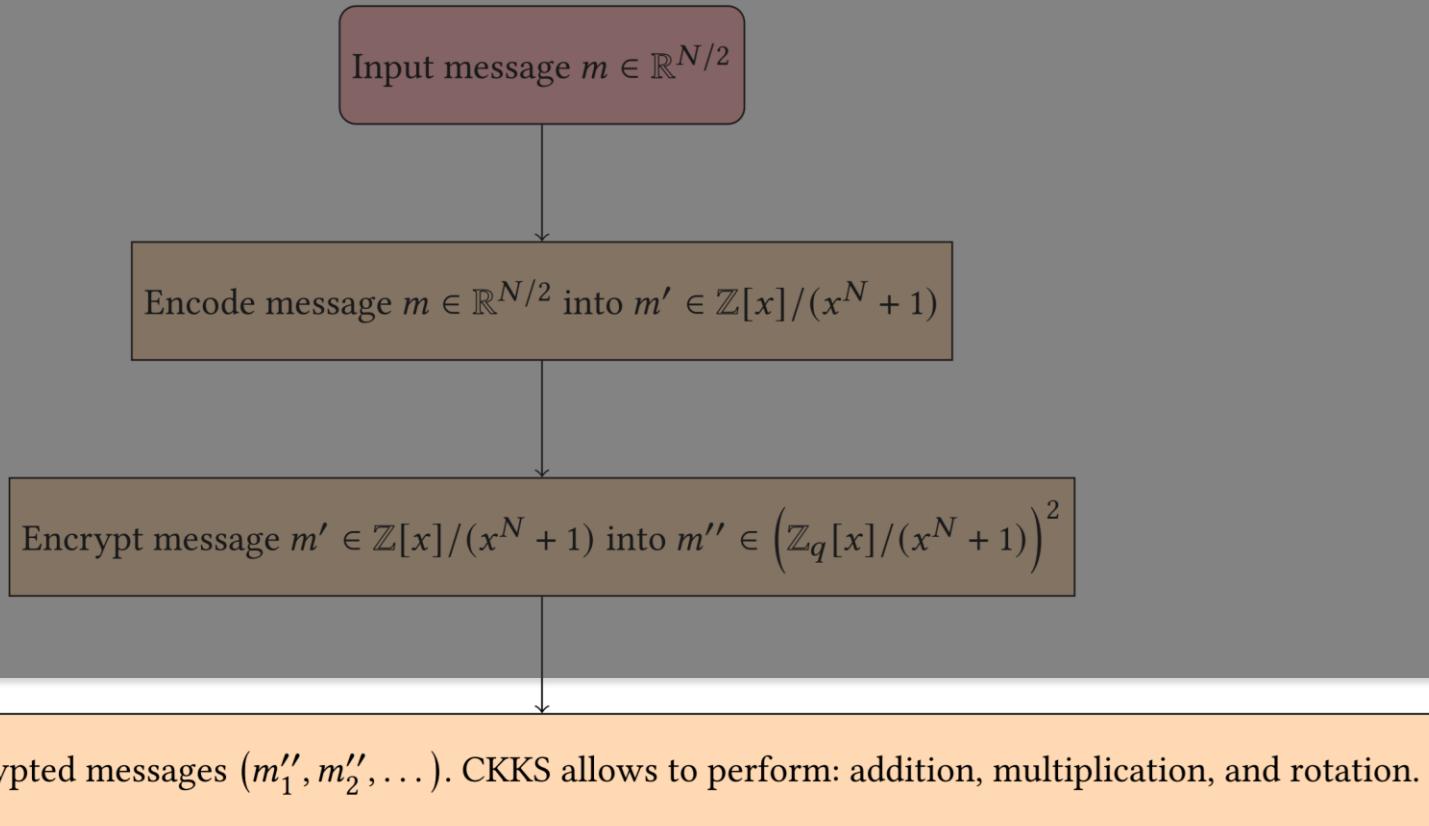
All keys in CKKS are polynomials from  $\mathbb{Z}_q[x]/(x^N + 1)$

The key size is  $Q \cdot N$  bits

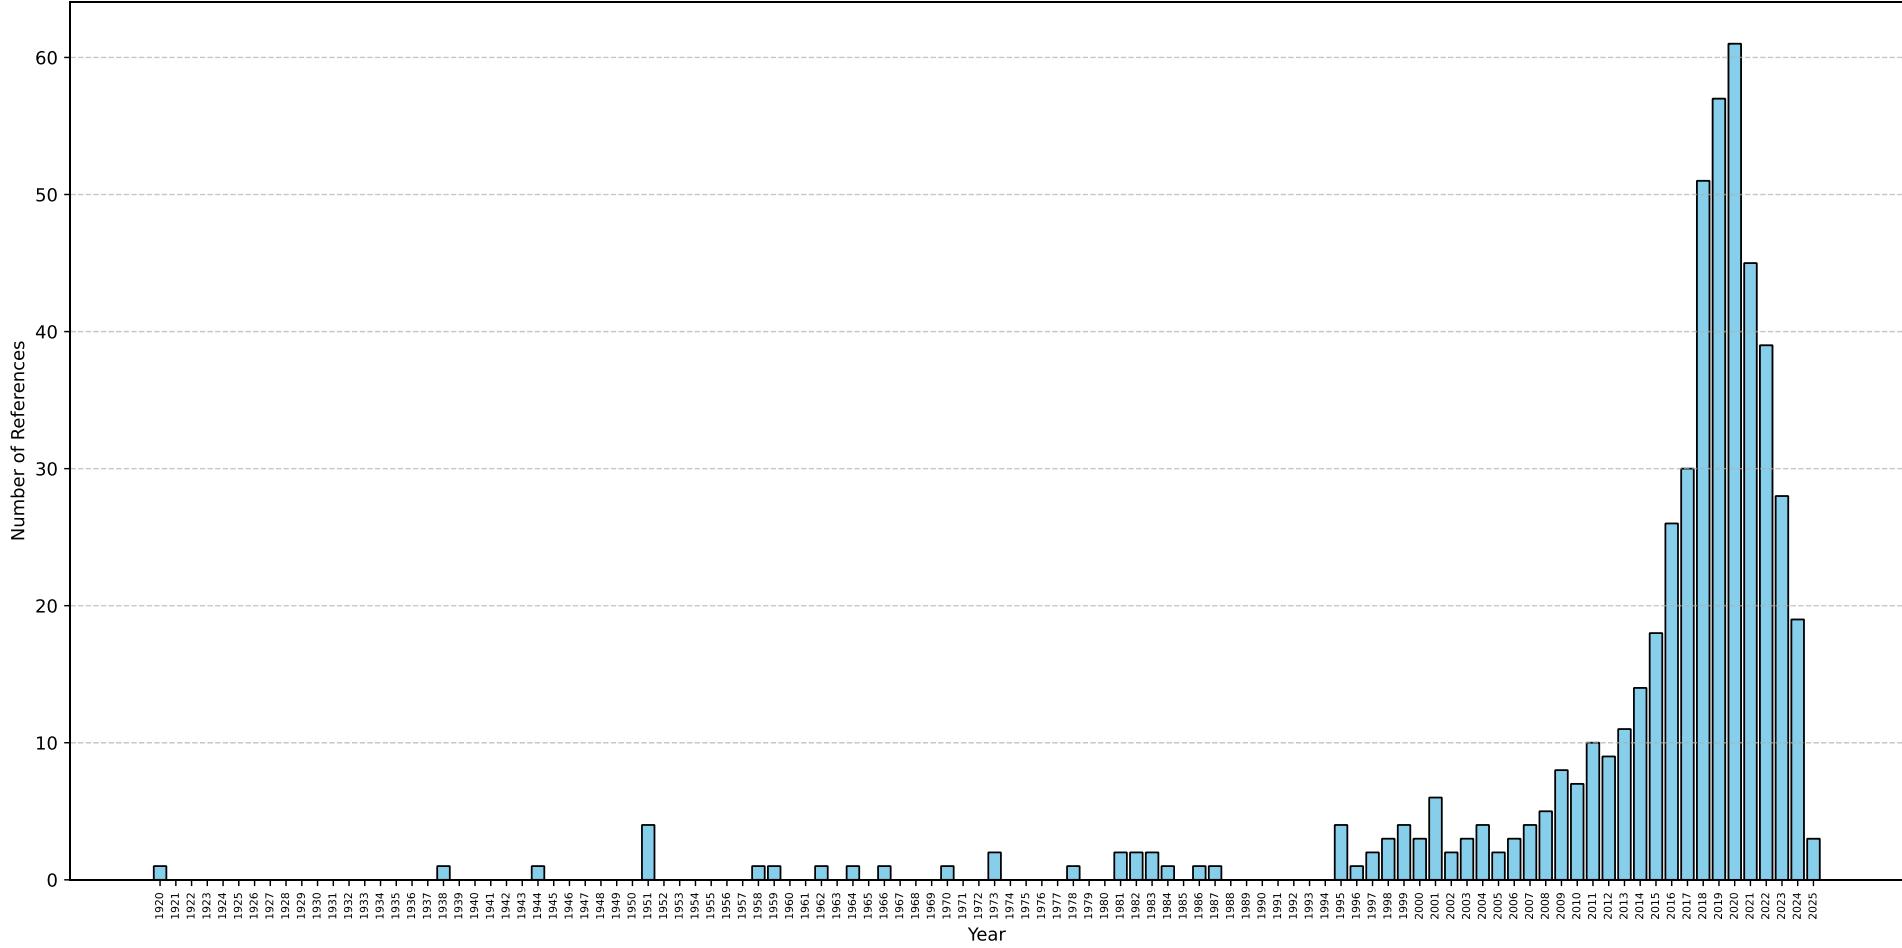
otation.

# Ch4: A Day of Life for Message with CKKS Scheme #54

• • •



## References by Year



# Ch3: EF21 Reloaded

#55

Original  $n = 4$  clients



$$L_1 = 1$$



$$L_2 = 1$$



$$L_3 = 1$$



$$L_4 = 1000$$

Original  $n = 4$  clients

## Example



$$L_1 = 1$$

$$L_{\text{AM}} = \frac{1}{4} (1 + 1 + 1 + 100) = 25.75$$



$$L_2 = 1$$



$$L_3 = 1$$

$$L_{\text{QM}} \approx \sqrt{\frac{1}{4} (1 + 1 + 1 + 100 \cdot 100)} = 15.73$$



$$L_4 = 1000$$

# Ch3: EF21 Reloaded (Approach 1)

#55

Original  $n = 4$  clients



$$L_1 = 1 \quad \lceil \frac{L_1}{L_{AM}} \rceil = \lceil \frac{1}{25.75} \rceil = 1$$



$$L_2 = 1 \quad \lceil \frac{L_2}{L_{AM}} \rceil = \lceil \frac{1}{25.75} \rceil = 1$$



$$L_3 = 1 \quad \lceil \frac{L_3}{L_{AM}} \rceil = \lceil \frac{1}{25.75} \rceil = 1$$



$$L_4 = 1000 \quad \lceil \frac{L_4}{L_{AM}} \rceil = \lceil \frac{1000}{25.75} \rceil = 39$$

Cloned  $N = 42$  clients



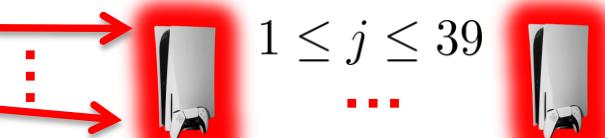
$$L_{11} = 1$$



$$L_{21} = 1$$



$$L_{31} = 1$$



Algorithm/Setting	Terminate condition	Iterations
Gradient Descent / $f(x)$ is strongly convex. (generalizable via adding regularizers with cheap proximal operator)	$ x^k - x^* ^2 \leq \varepsilon  x^0 - x^* ^2$ $\Rightarrow \text{if } \nabla f(x^*) = 0 \text{ then }  f(x^k) - f(x^*)  \leq \frac{L}{2}  x^k - x^* ^2$	$k \geq \frac{L}{\mu} \log\left(\frac{1}{\varepsilon}\right)$
Stochastic Gradient Descent / $f(x)$ is strongly convex / $g(x)$ such that $E[g(x) x] = \nabla f(x)$ / $g(x)$ such that $E[ \nabla g(x) - \nabla f(x^*) ^2  x] \leq 2AD_f(x, x^*) + C$ $D_f(x, x^*) = f(x) - f(x^*) - \langle \nabla f(x^*), x - x^* \rangle$	$E[ x^k - x^* ^2] \leq \varepsilon  x^0 - x^* ^2$ $\Rightarrow \text{if } \nabla f(x^*) = 0 \text{ then }  f(x^k) - f(x^*)  \leq \frac{L}{2}  x^k - x^* ^2$	$k \geq O\left(\frac{1}{\varepsilon} \cdot \log\left(\frac{1}{\varepsilon}\right)\right)$
Gradient Descent / $f(x)$ is convex.	$f(x^k) - f(x^*) \leq \varepsilon  x^0 - x^* ^2$	$k \geq \frac{1}{2\alpha\varepsilon}, \alpha \in \left(0, \frac{1}{L}\right]$
Accelerate Gradient Descent / $f(x)$ is convex.	$f(x^k) - f(x^*) \leq \varepsilon  x^0 - x^* ^2$	$k \geq 1 + \sqrt{\frac{2}{\alpha\varepsilon}}, \alpha = \frac{1}{L}$ <span style="float: right;">(Optimal)</span>
Stochastic Subgradient Descent / $f(x)$ is convex / $g(x)$ is unbiased / easy prove: $g(x)$ is bounded by $G$ , start iterate $x^0$ has an upper bound distance $R$ to $x^*$	$E[f(x^k) - f(x^*)] \leq \varepsilon$	$k \geq O\left(\frac{1}{\varepsilon^2}\right)$ with optimal $a_k = \frac{R/G}{\sqrt{k}}$ <span style="float: right;">(Optimal)</span>
Stochastic Gradient Descent / $f(x)$ is convex / $g(x)$ is unbiased / $g(x)$ satisfy sigma-k assumption	$E[f(x^k) - f(x^*)] \leq \varepsilon$	$k \geq O\left(\frac{1}{\varepsilon}\right)$ to neighborhood. $k \geq O\left(\frac{1}{\varepsilon^2}\right)$ exactly convergence to a solution. <span style="float: right;">(Optimal)</span>
Gradient Descent / $f(x)$ is non-convex, but smooth	$ \nabla f(x^k)  \leq \varepsilon$	$k \geq O\left(\frac{1}{\varepsilon^2}\right)$
Stochastic Gradient Descent / $f(x)$ is non-convex, but smooth.	$E \nabla f(x^k)  \leq \varepsilon$	From $k \geq O\left(\frac{1}{\varepsilon^2}\right)$ to $k \geq O\left(\frac{1}{\varepsilon^4}\right)$ depending on assumptions
Optimal SGD for case when $f(x)$ is finite sum of $n$ functions / $f(x)$ is non-convex, but smooth (e.g. PAGE) PAGE reduces to GD if $p=1$ or $\tau = n$	$ \nabla f(x^k)  \leq \varepsilon$	$k \geq O\left(\frac{\sqrt{n}}{\varepsilon^2}\right)$ <span style="float: right;">(Optimal)</span>